

Characterize the Performance of cache and TLB

袁浩程
圣贤

2021/11/22

Problem Statement

Measure cache and TLB

- I have a computer
- I want to know the cpu's....



	<i>Cache Parameters</i>	<i>TLB Parameters</i>
<i>cache size</i>	?	?
<i>associativity</i>	?	?
<i>line size</i>	?	?
<i>miss penalty</i>	?	?

Experimental Methodology

Basic idea

- a: associativity
- b: line size
- D: cache size
- M: miss penalty
- N: array size
- S: stride

- Compute many times a simple floating-point function on a subset of elements taken from an array.
- This subset is given by
$$1, s + 1, 2s + 1, \dots, N - s + 1 (s < N/2)$$
- Each iteration the cache gets a read request, time differs when miss does/doesn't happen.

Experimental Methodology

Four regimes

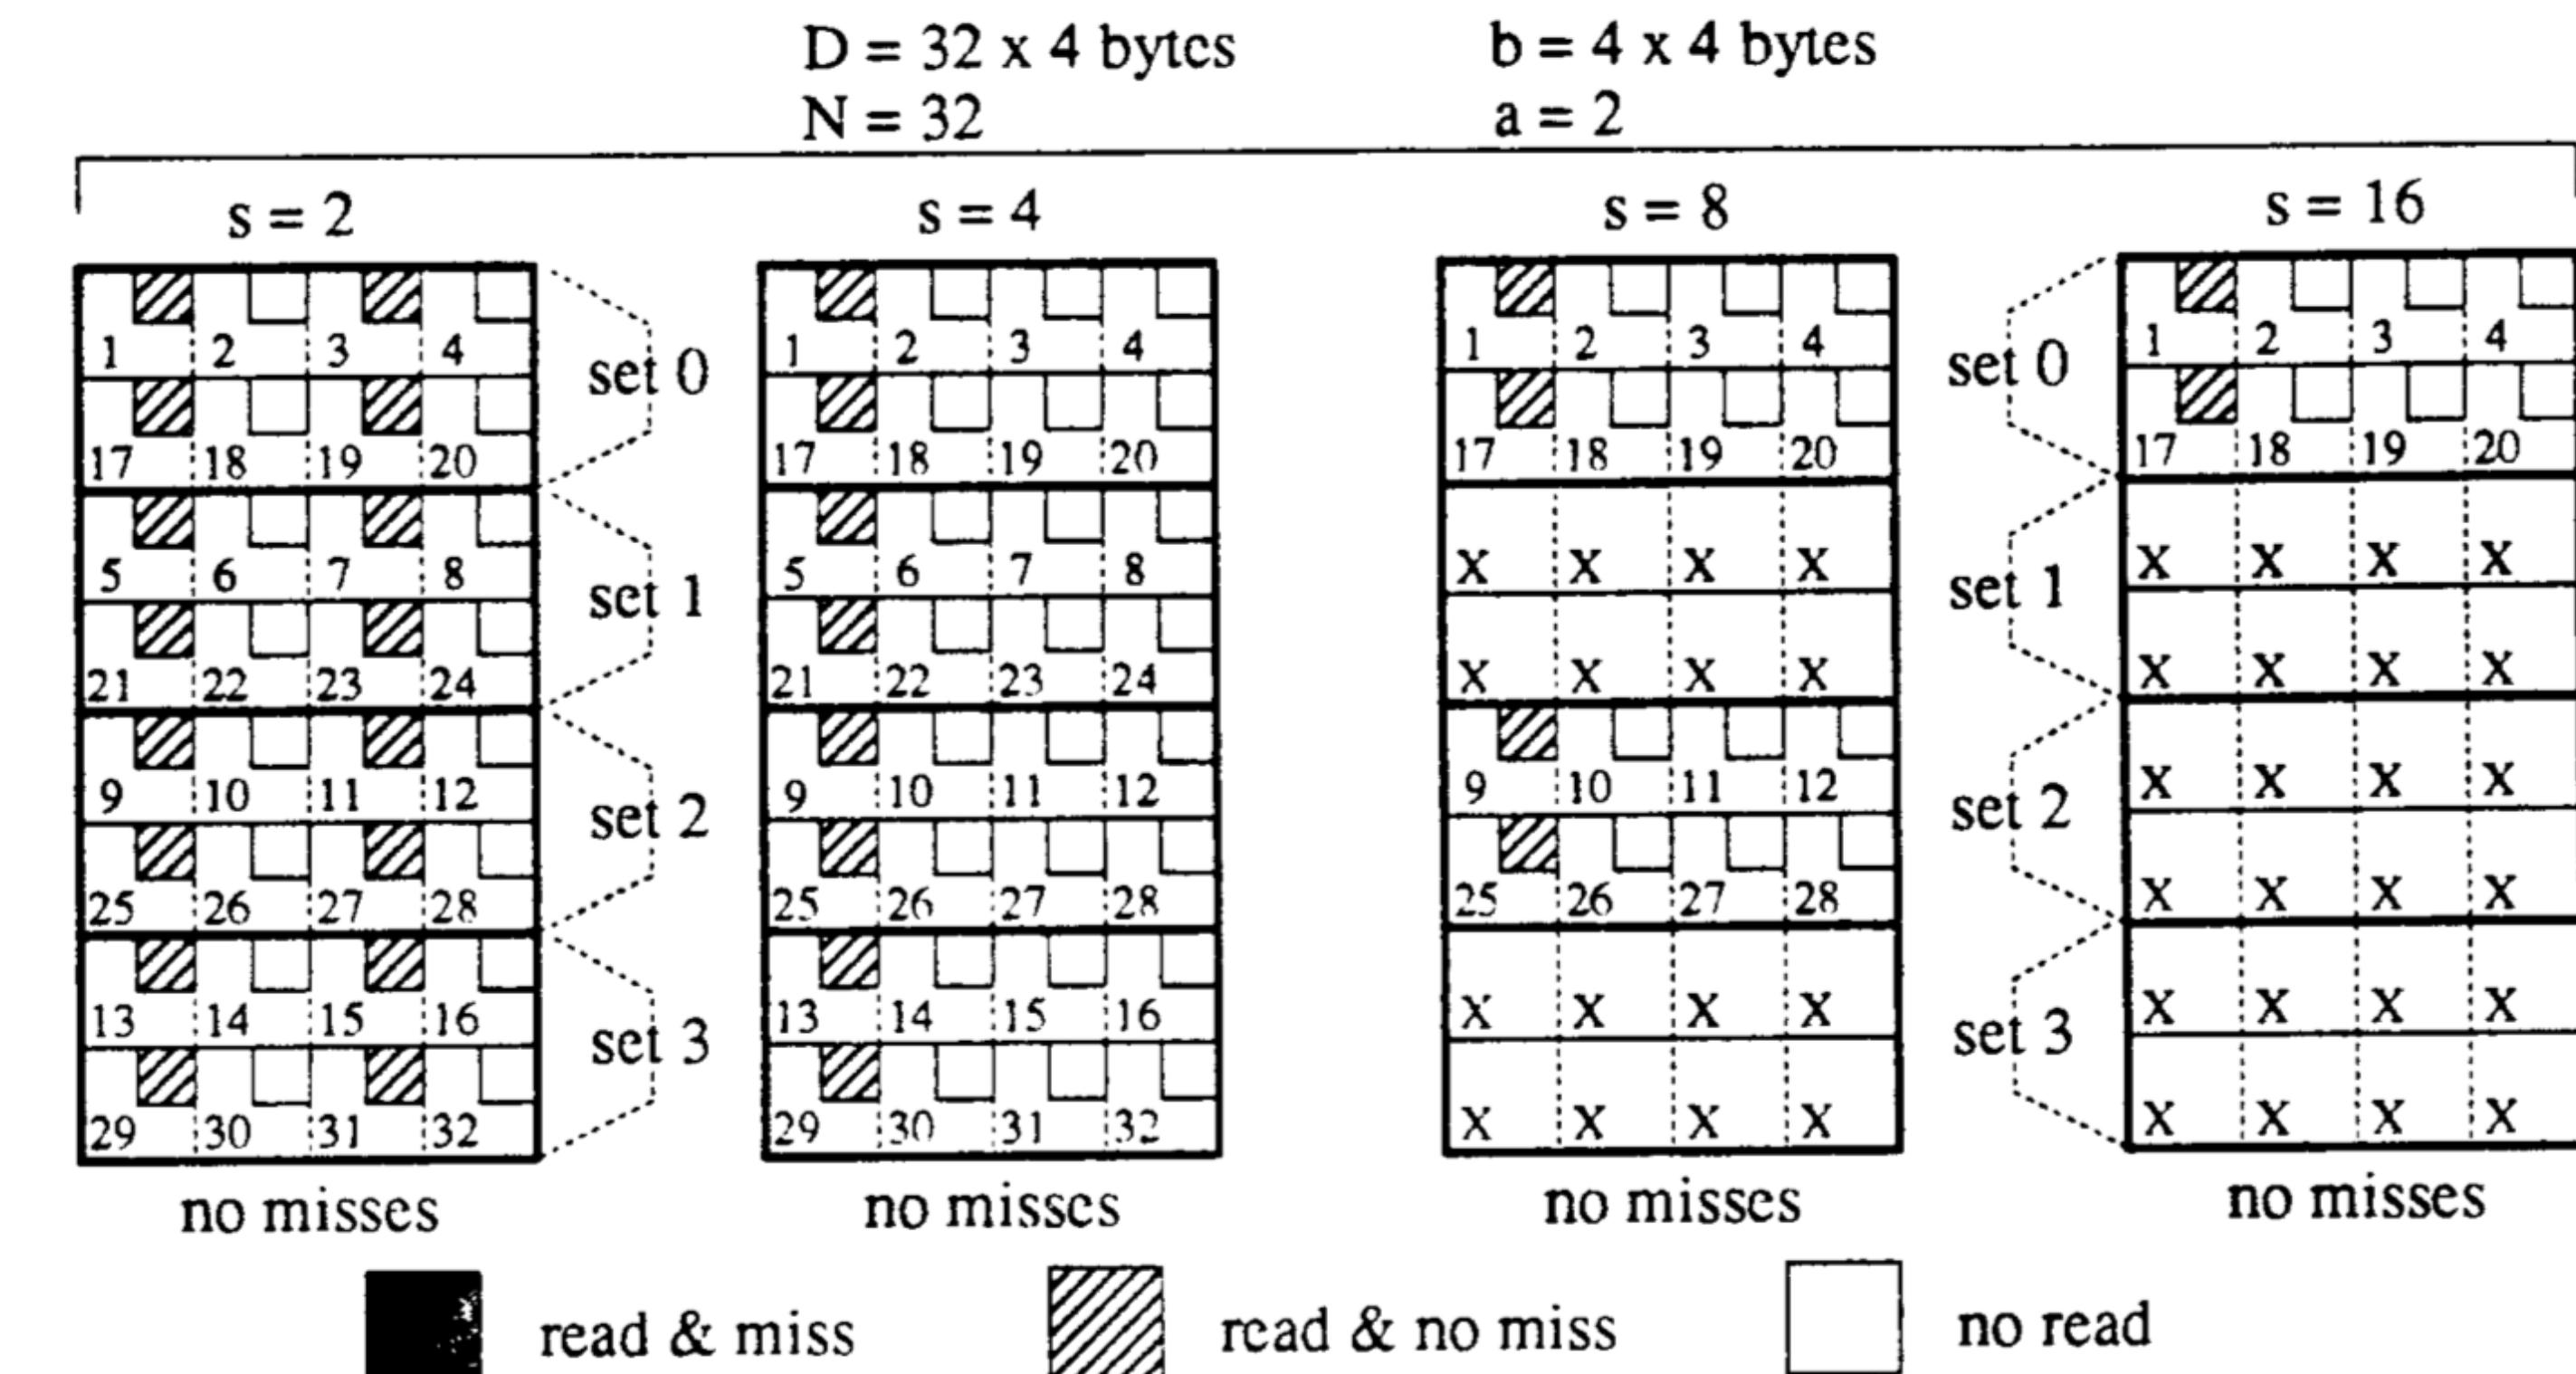
- a: associativity
- b: line size
- D: cache size
- M: miss penalty
- N: array size
- S: stride

<i>Regime</i>	<i>Size of Array</i>	<i>Stride</i>	<i>Frequency of Misses</i>	<i>Time per Iteration</i>
1	$1 < N < D$	$1 < s < N/2$	no misses	$T_{\text{no_miss}}$
2.a	$D < N$	$1 \leq s < b$	one miss every b/s elements	$T_{\text{no_miss}} + M s/b$
2.b	$D < N$	$b \leq s < N/a$	one miss every element	$T_{\text{no_miss}} + M$
2.c	$D < N$	$N/a \leq s \leq N/2$	no misses	$T_{\text{no_miss}}$

Experimental Methodology

Four regimes

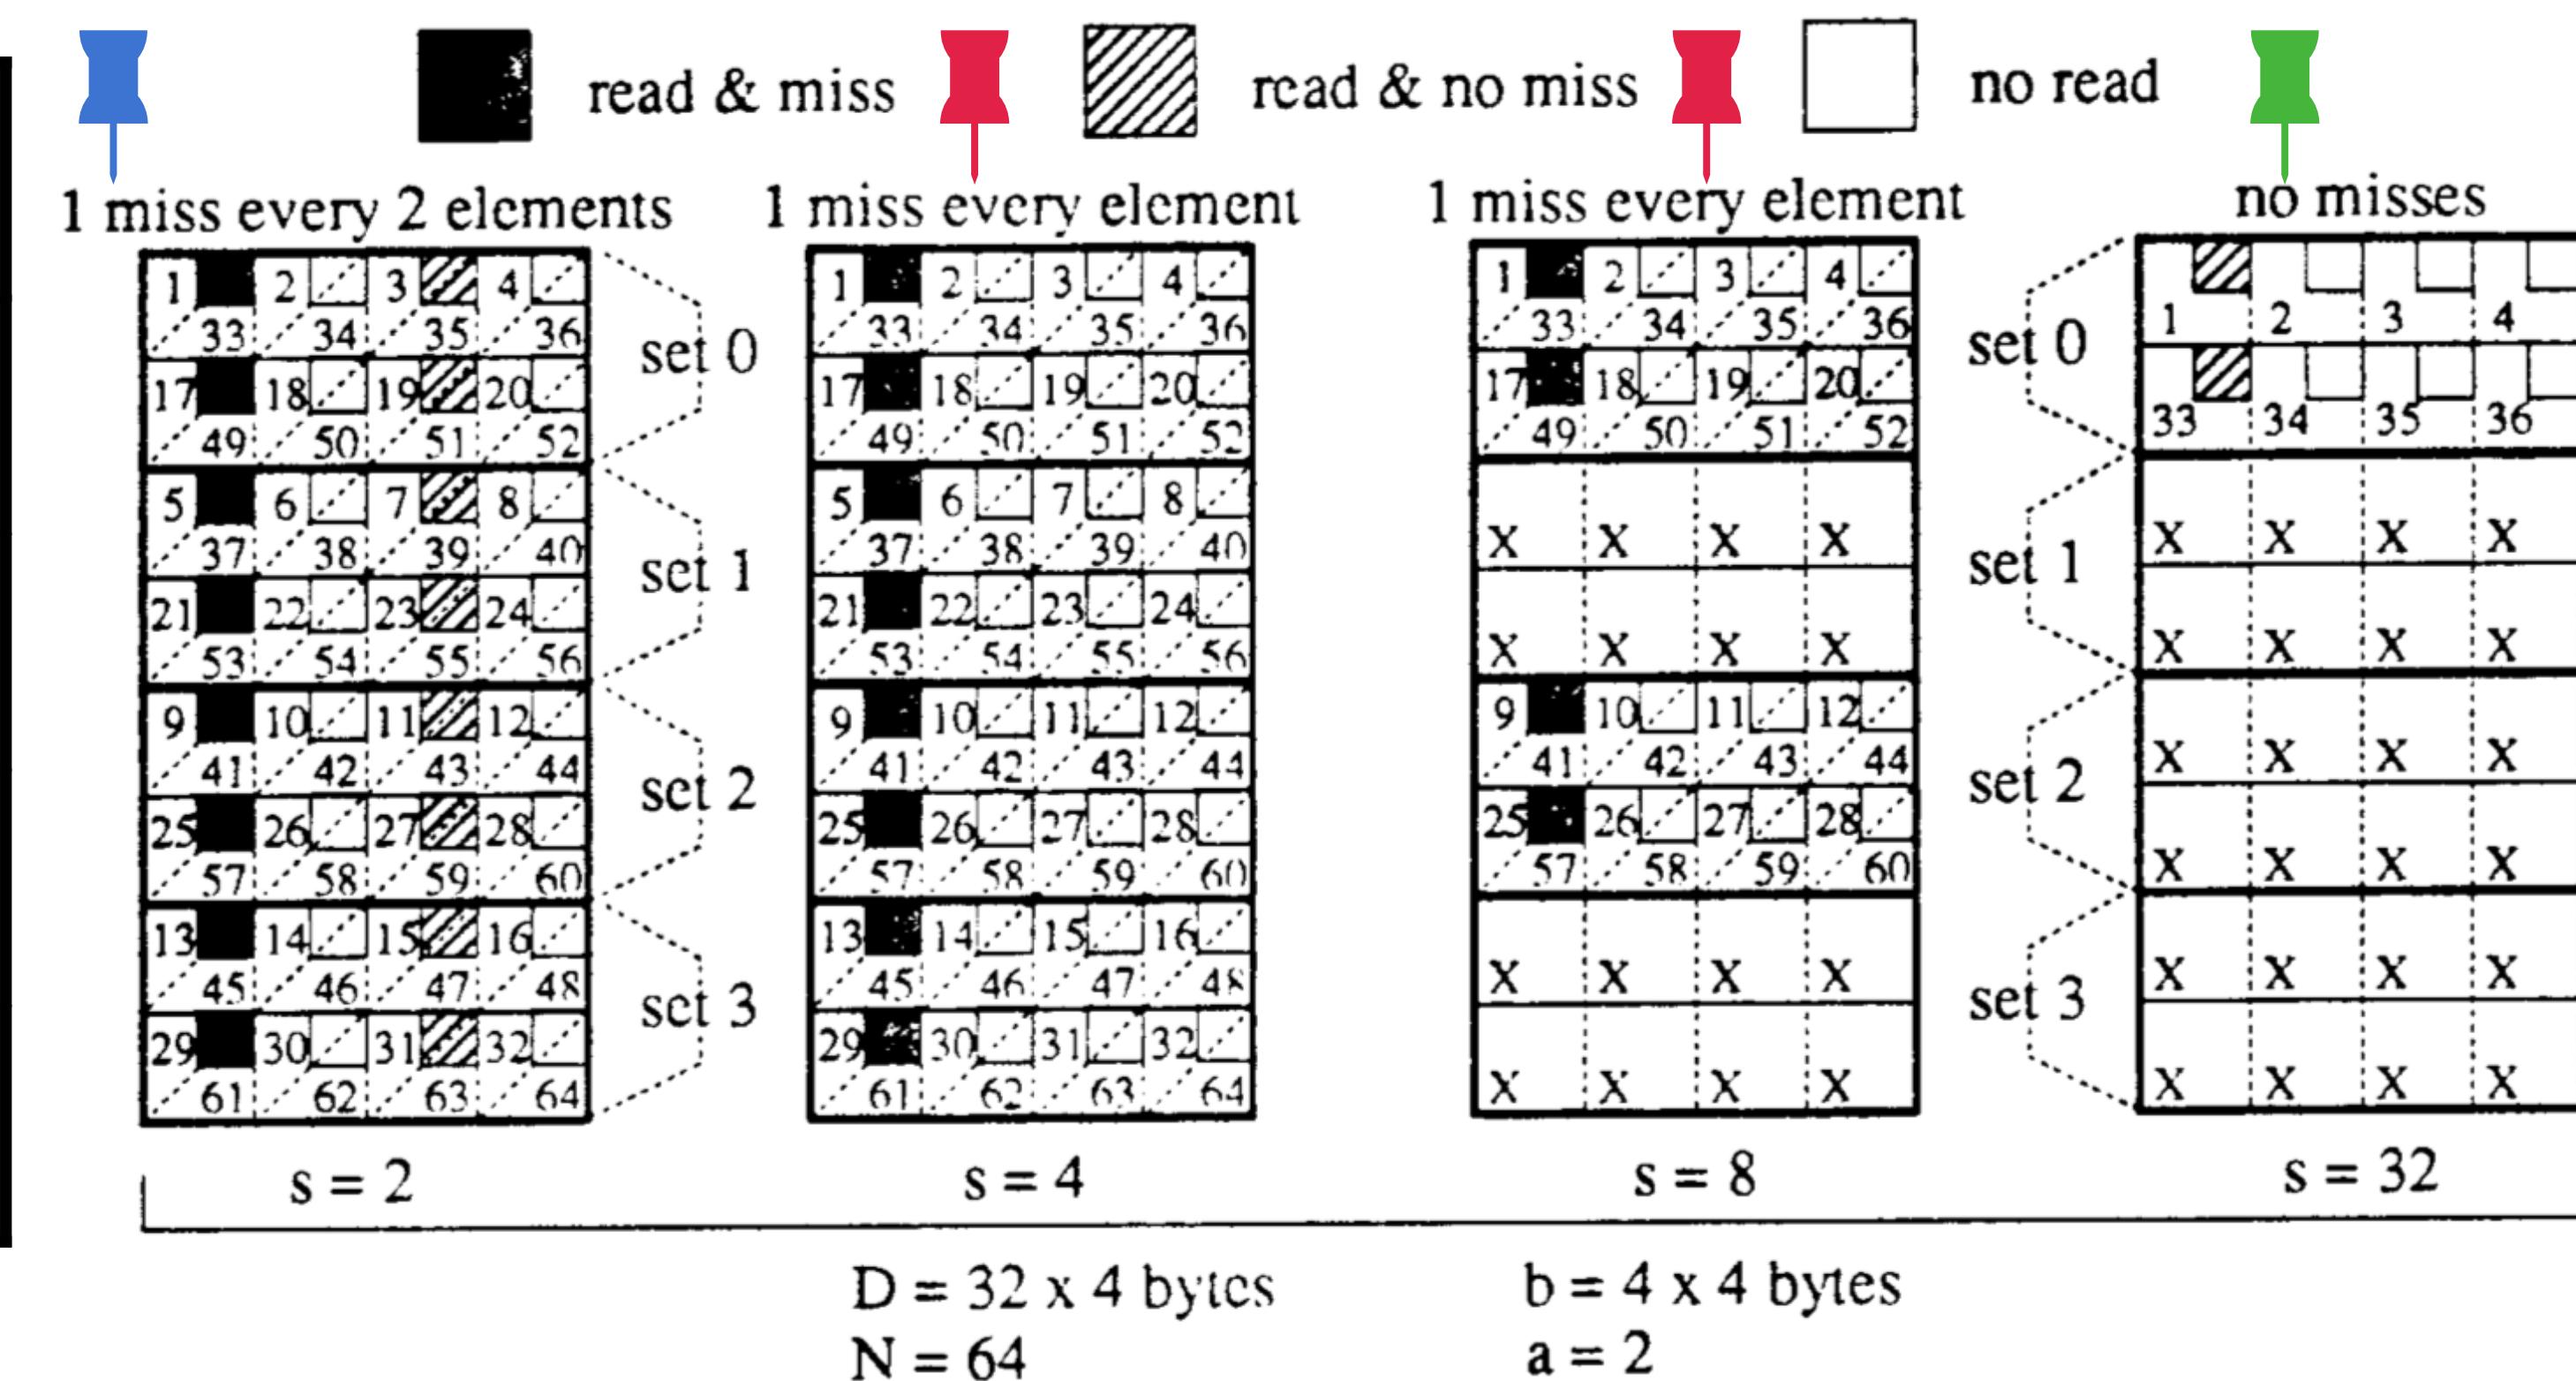
<i>Regime</i>	<i>Size of Array</i>	<i>Stride</i>	<i>Frequency of</i>	<i>Time per Iteration</i>
1	$1 < N < D$	$1 < s < N/2$	no	T
2.a	$D < N$	$1 \leq s < b$	one every b/a	$T + Ms/b$
2.b	$D < N$	$b \leq s < N/a$	one every	$T + M$
2.c	$D < N$	$N/a \leq s \leq N/2$	no	T



Experimental Methodology

Four regimes

<i>Regime</i>	<i>Size of Array</i>	<i>Stride</i>	<i>Frequency of</i>	<i>Time per Iteration</i>
1	$1 < N < D$	$1 < s < N/2$	no	T
2.a	$D < N$	$1 \leq s < b$	one per b/s	$T + Ms/b$
2.b	$D < N$	$b \leq s < N/a$	one	$T + M$
2.c	$D < N$	$N/a \leq s \leq N/2$	no	T

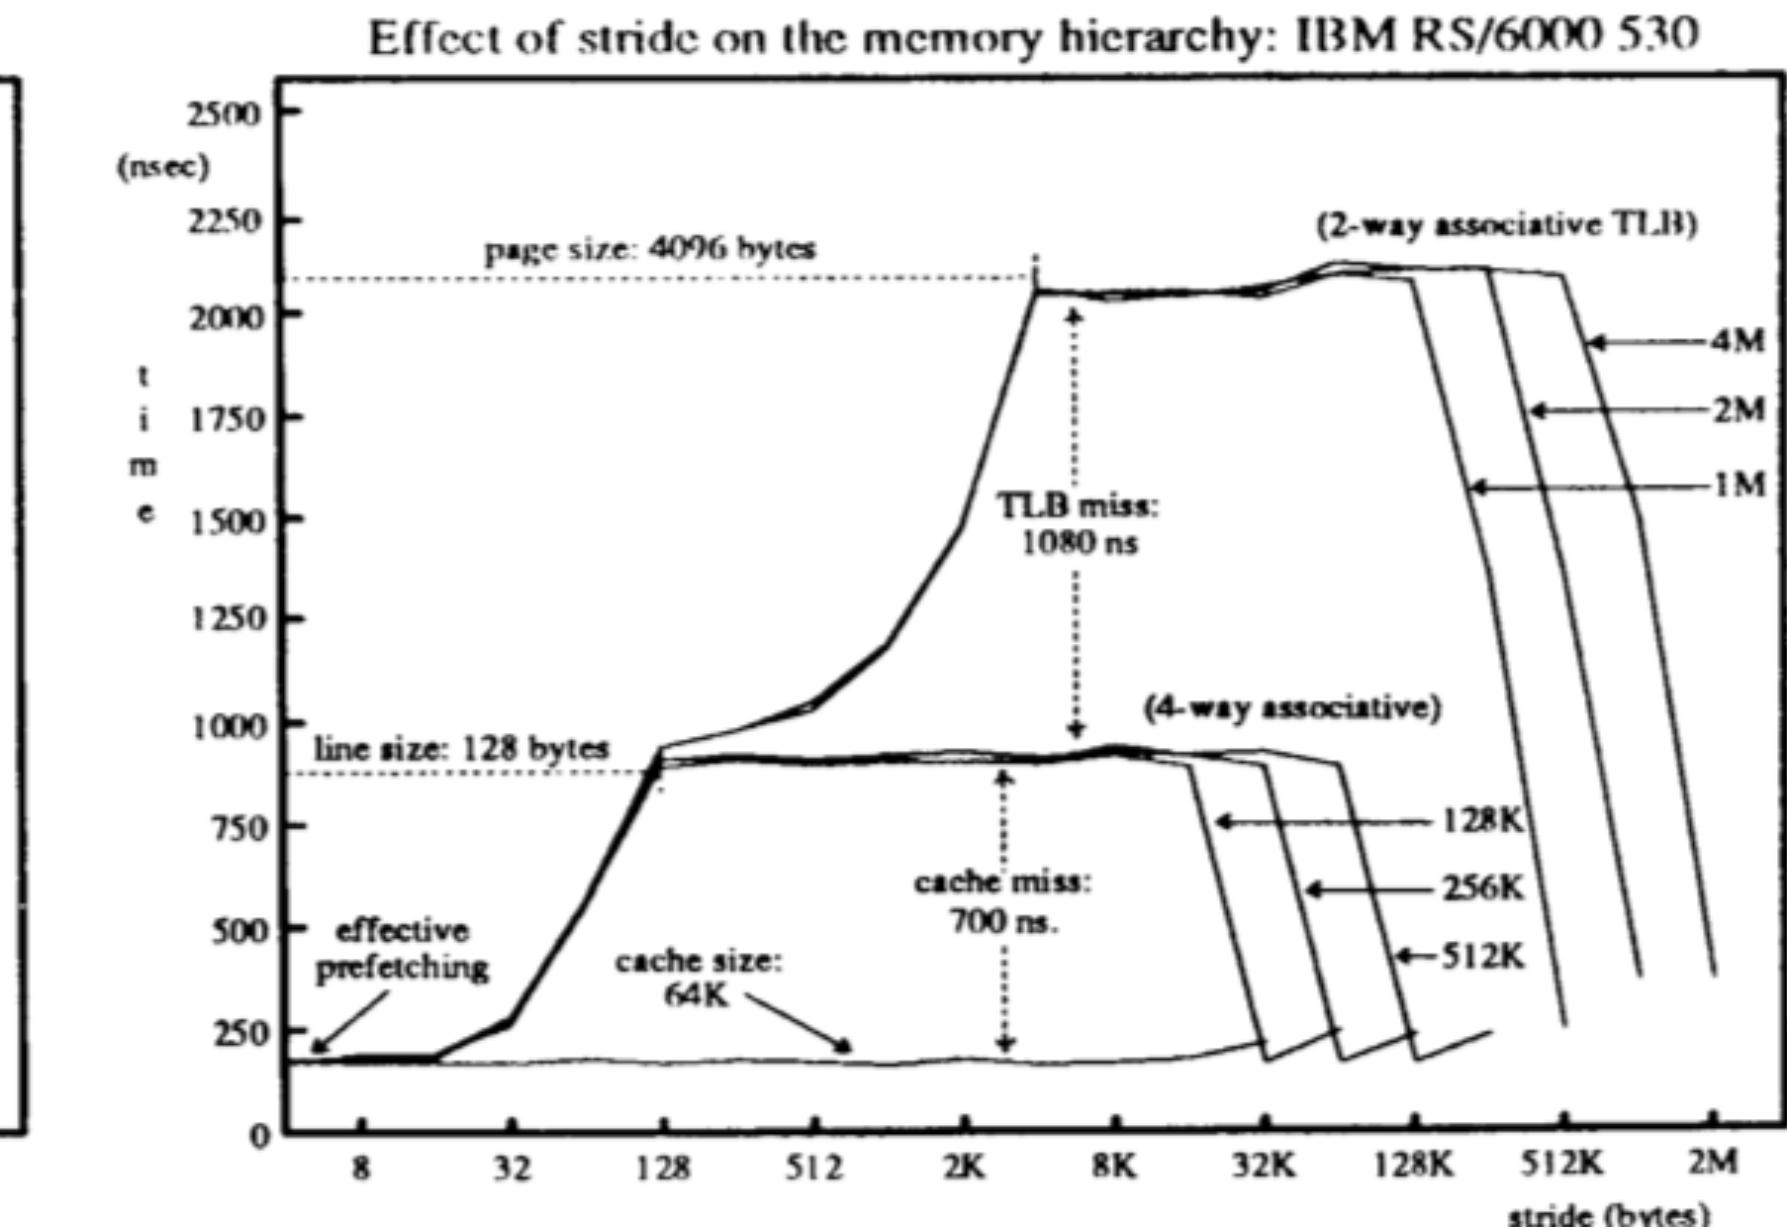
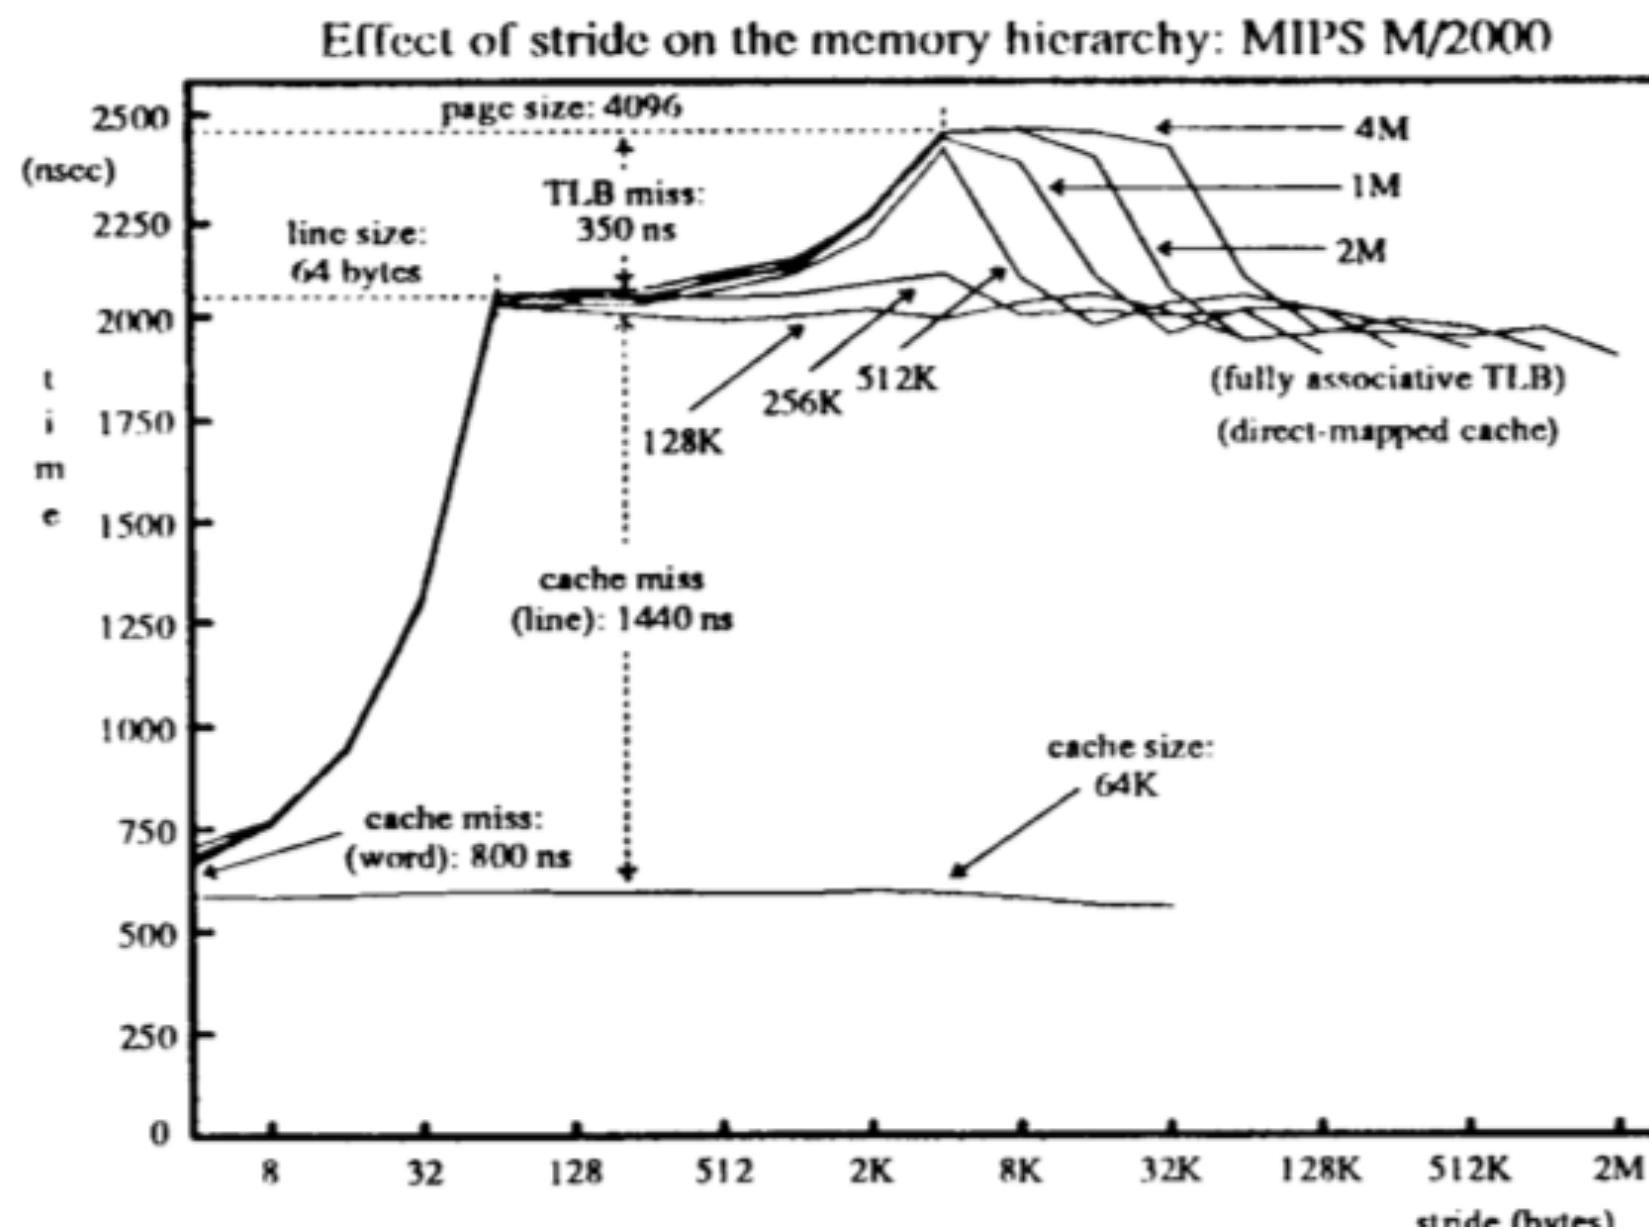


Experimental Methodology

My implementation

- C build the main part of the benchmark
- Shell change N & S automatically
- Matlab draw the graph

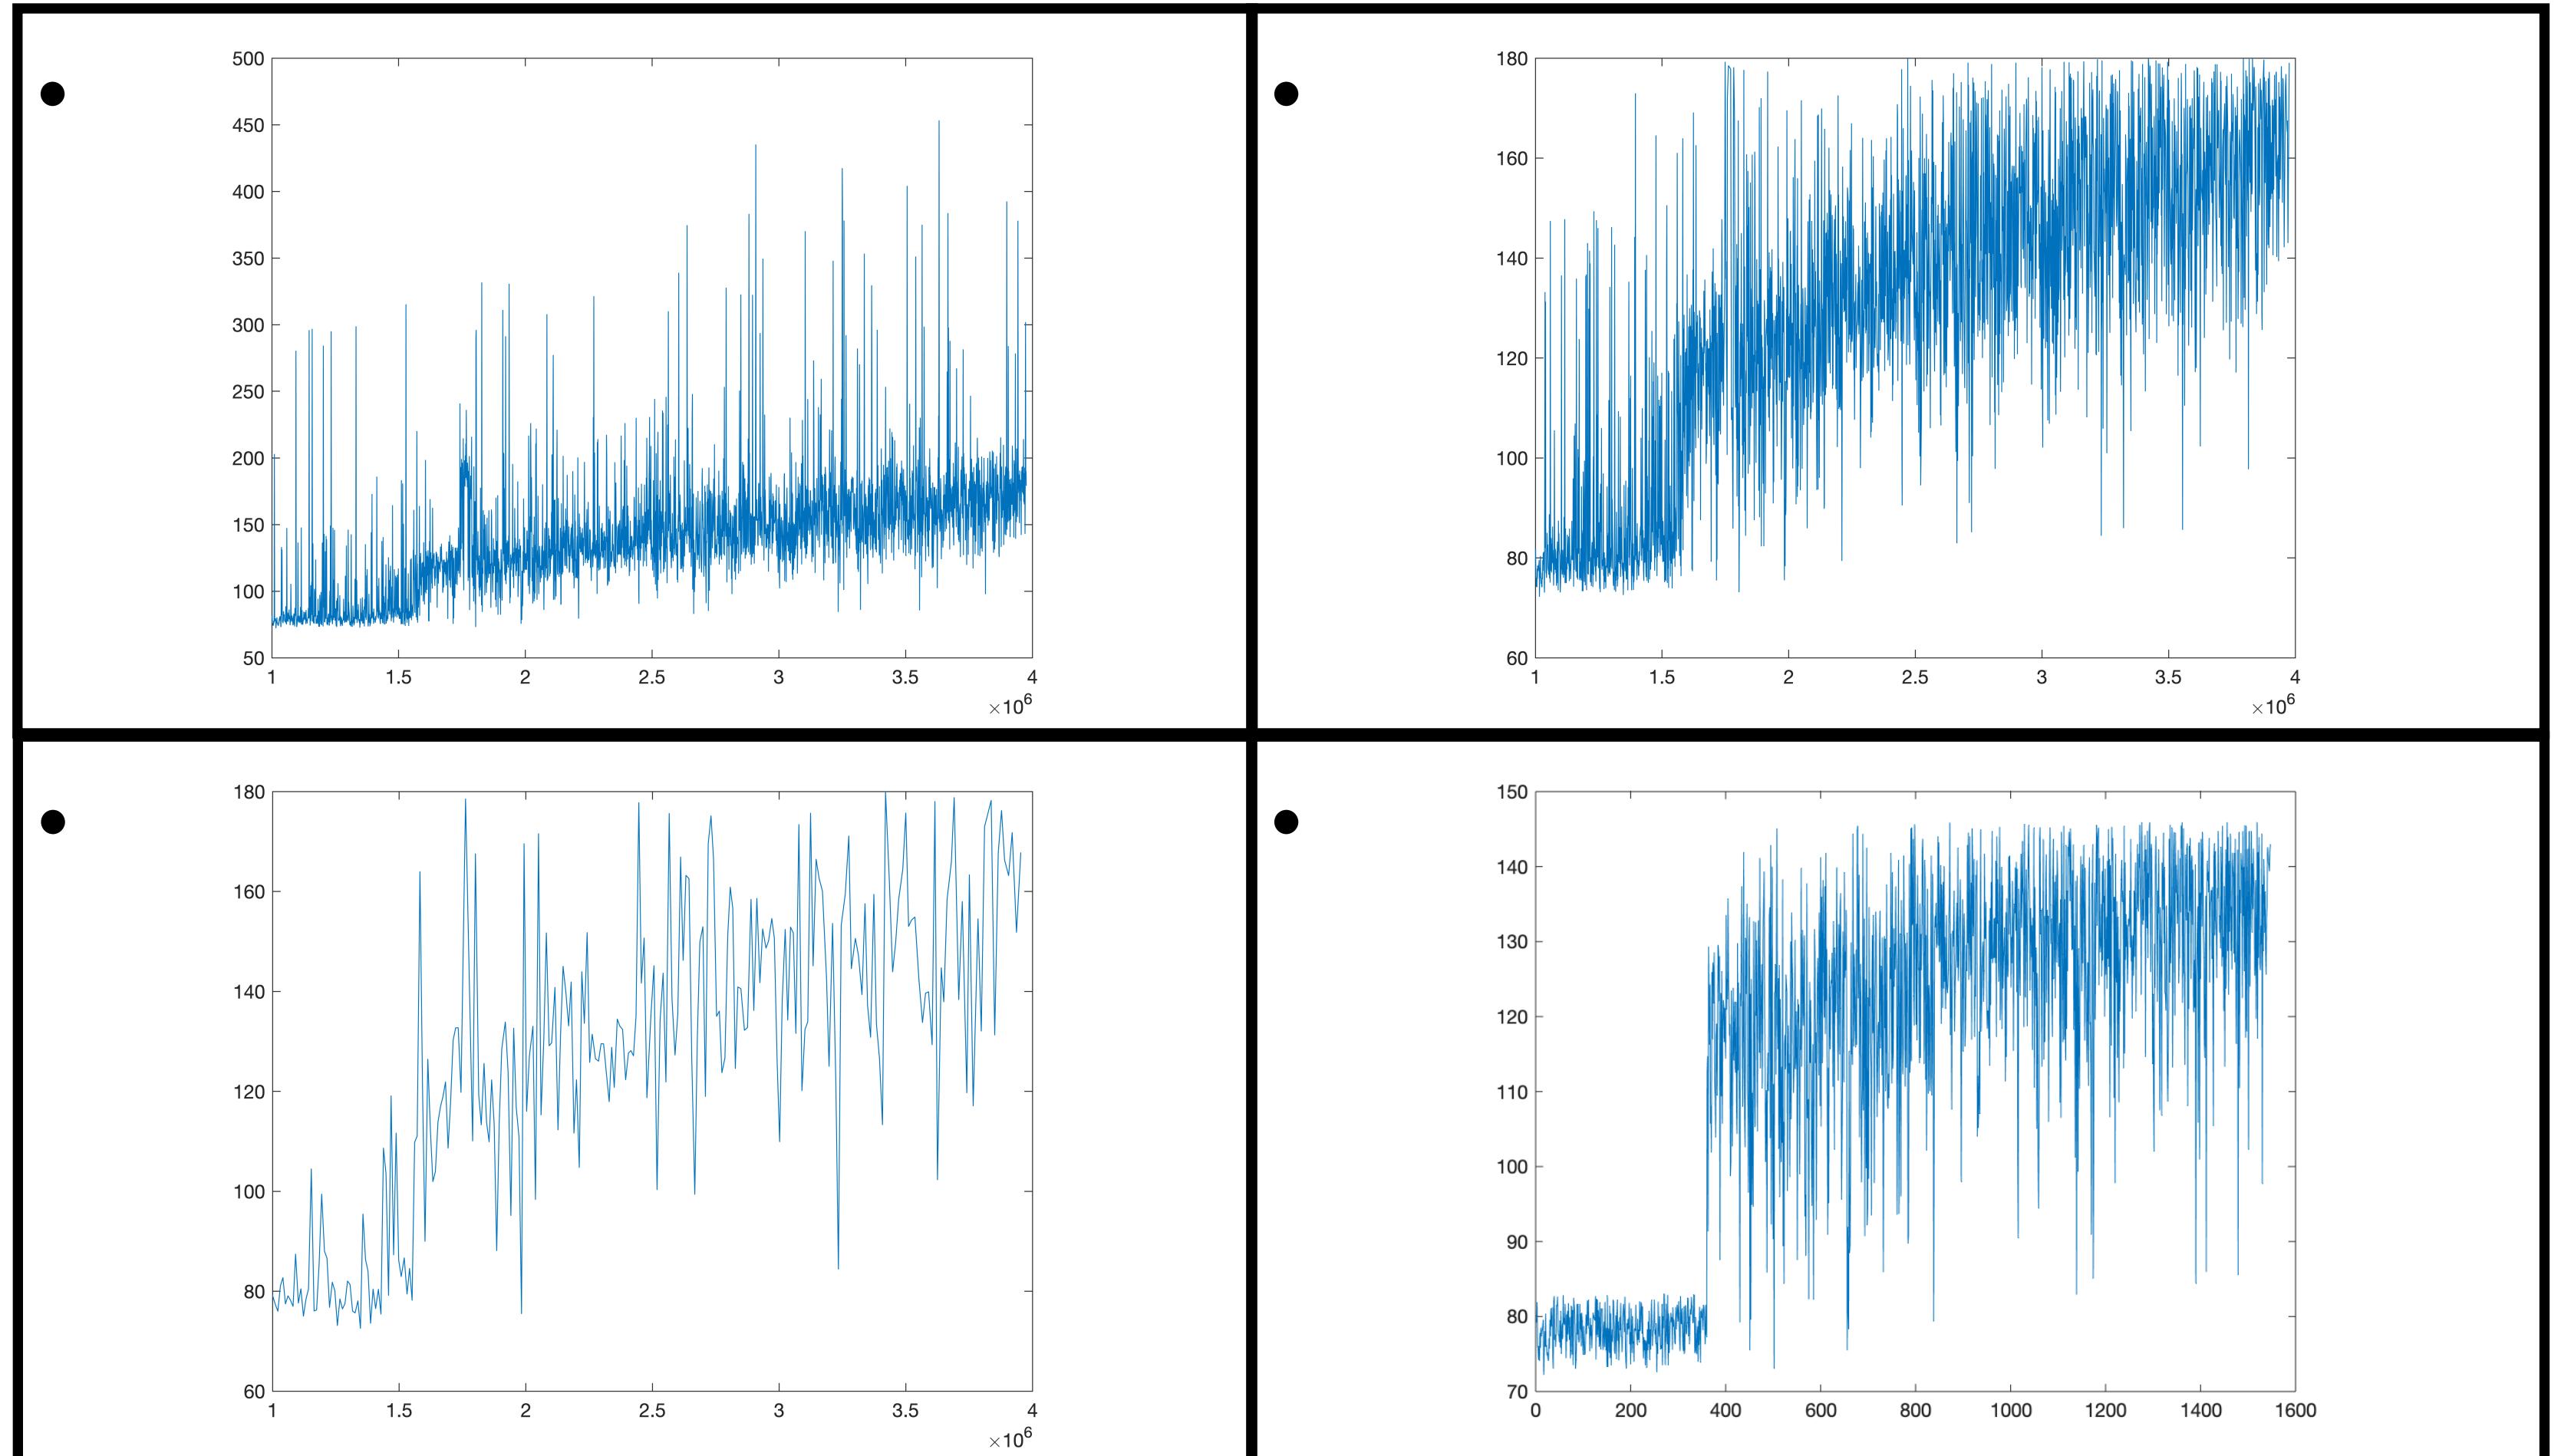
- Use benchmark_N to identify the cache size
- Use benchmark_S to get the result



Result

Benchmark_N

- benchmark_n.c
 - *build the main loop*
 - *count the time*
 - *write the result*
- benchmark_n.sh
 - *change N automatically*

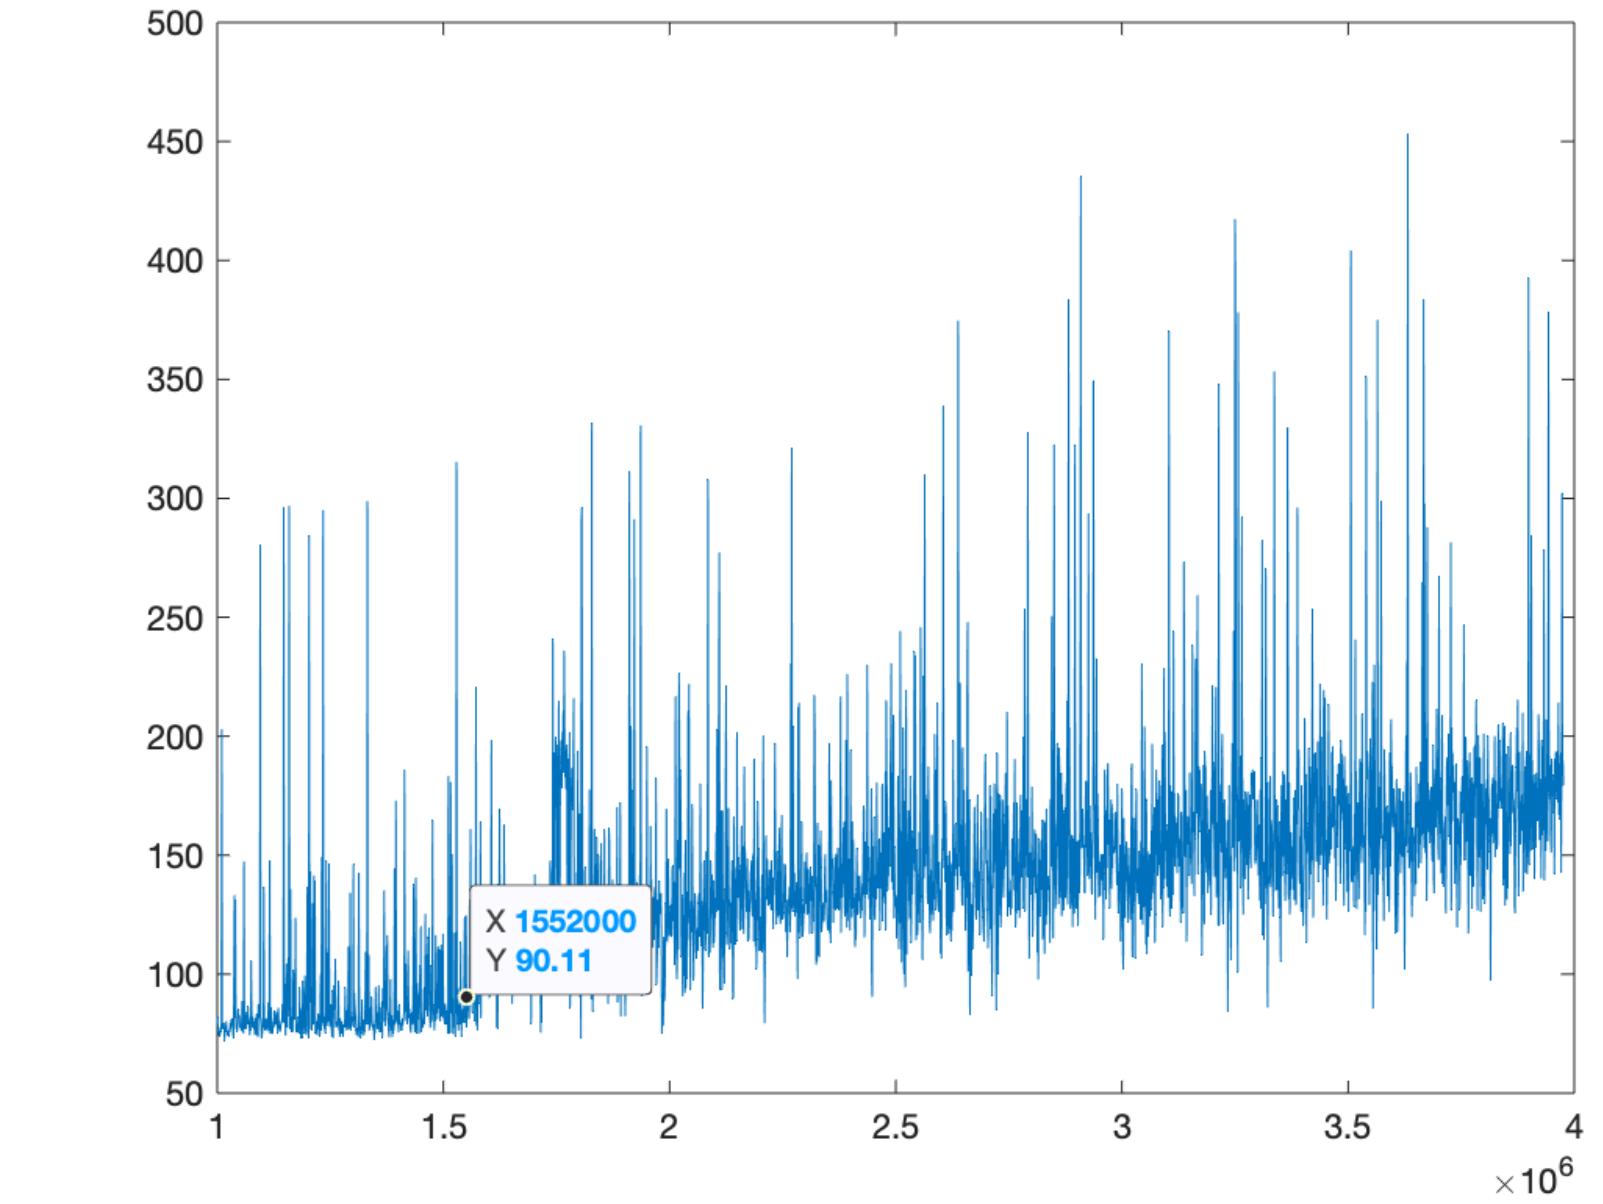
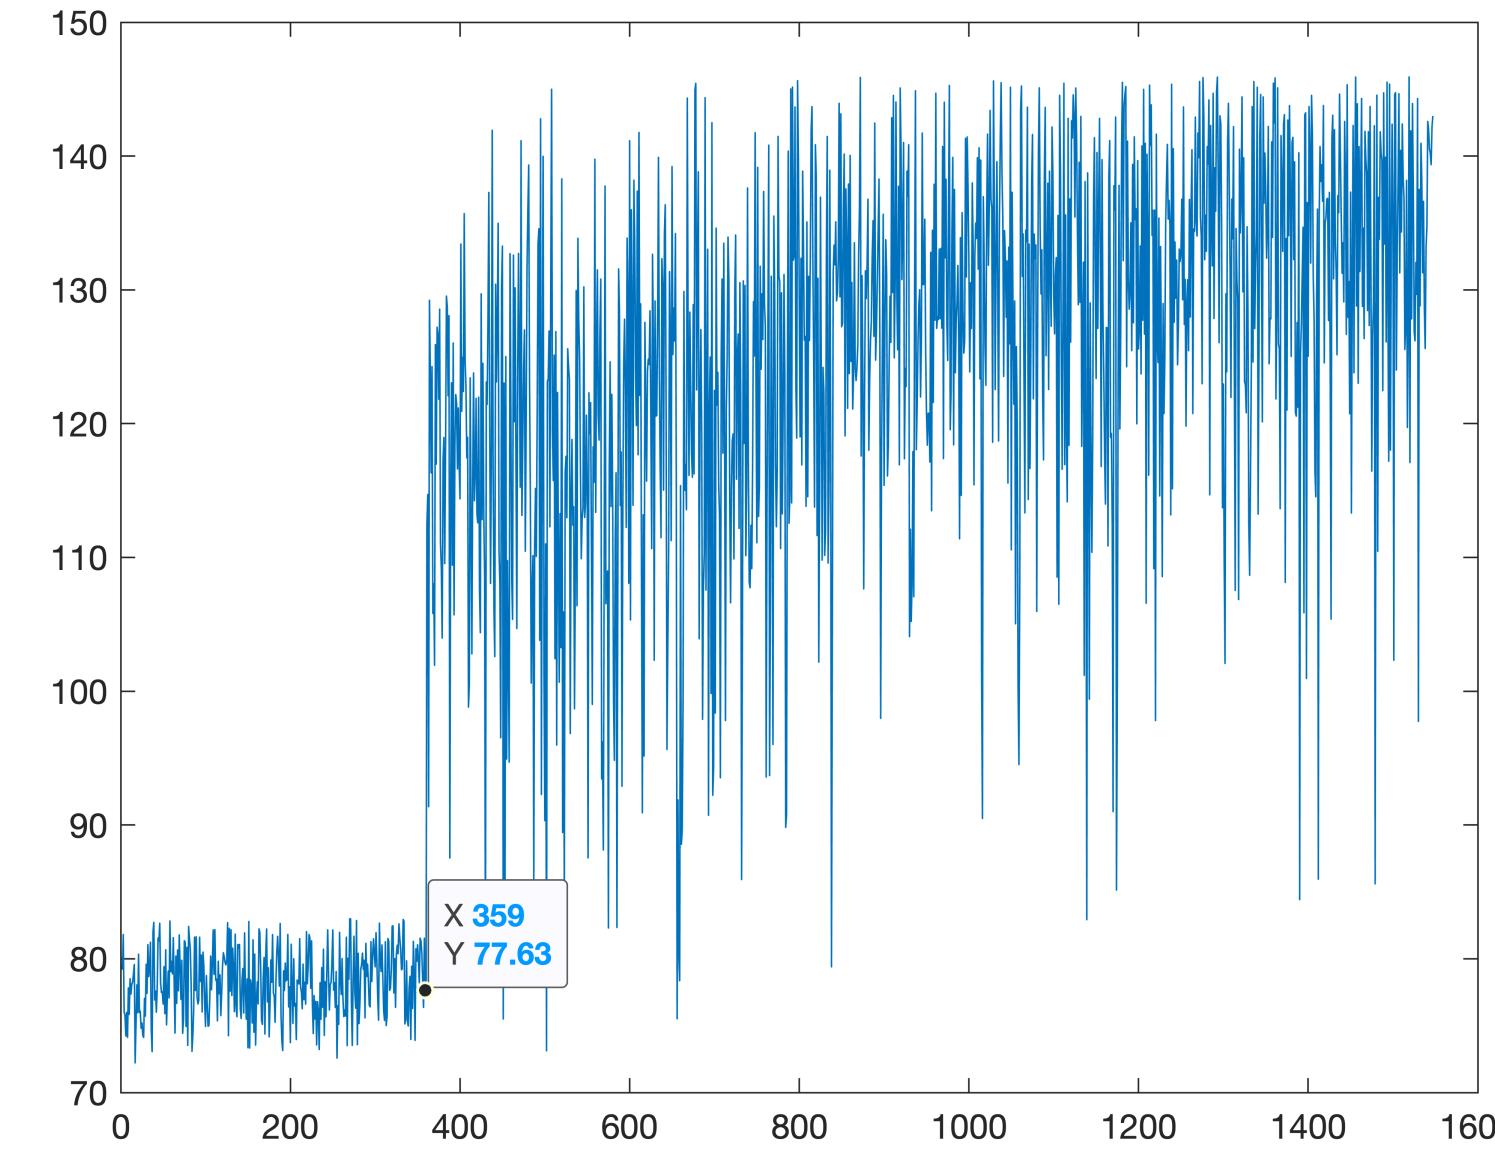


Result

Benchmark_N

- benchmark_n.c
 - *build the main loop*
 - *count the time*
 - *write the result*
- benchmark_n.sh
 - *change N automatically*

- 359 or 400 ?
- 1552000 4-bytes words = 6062KB

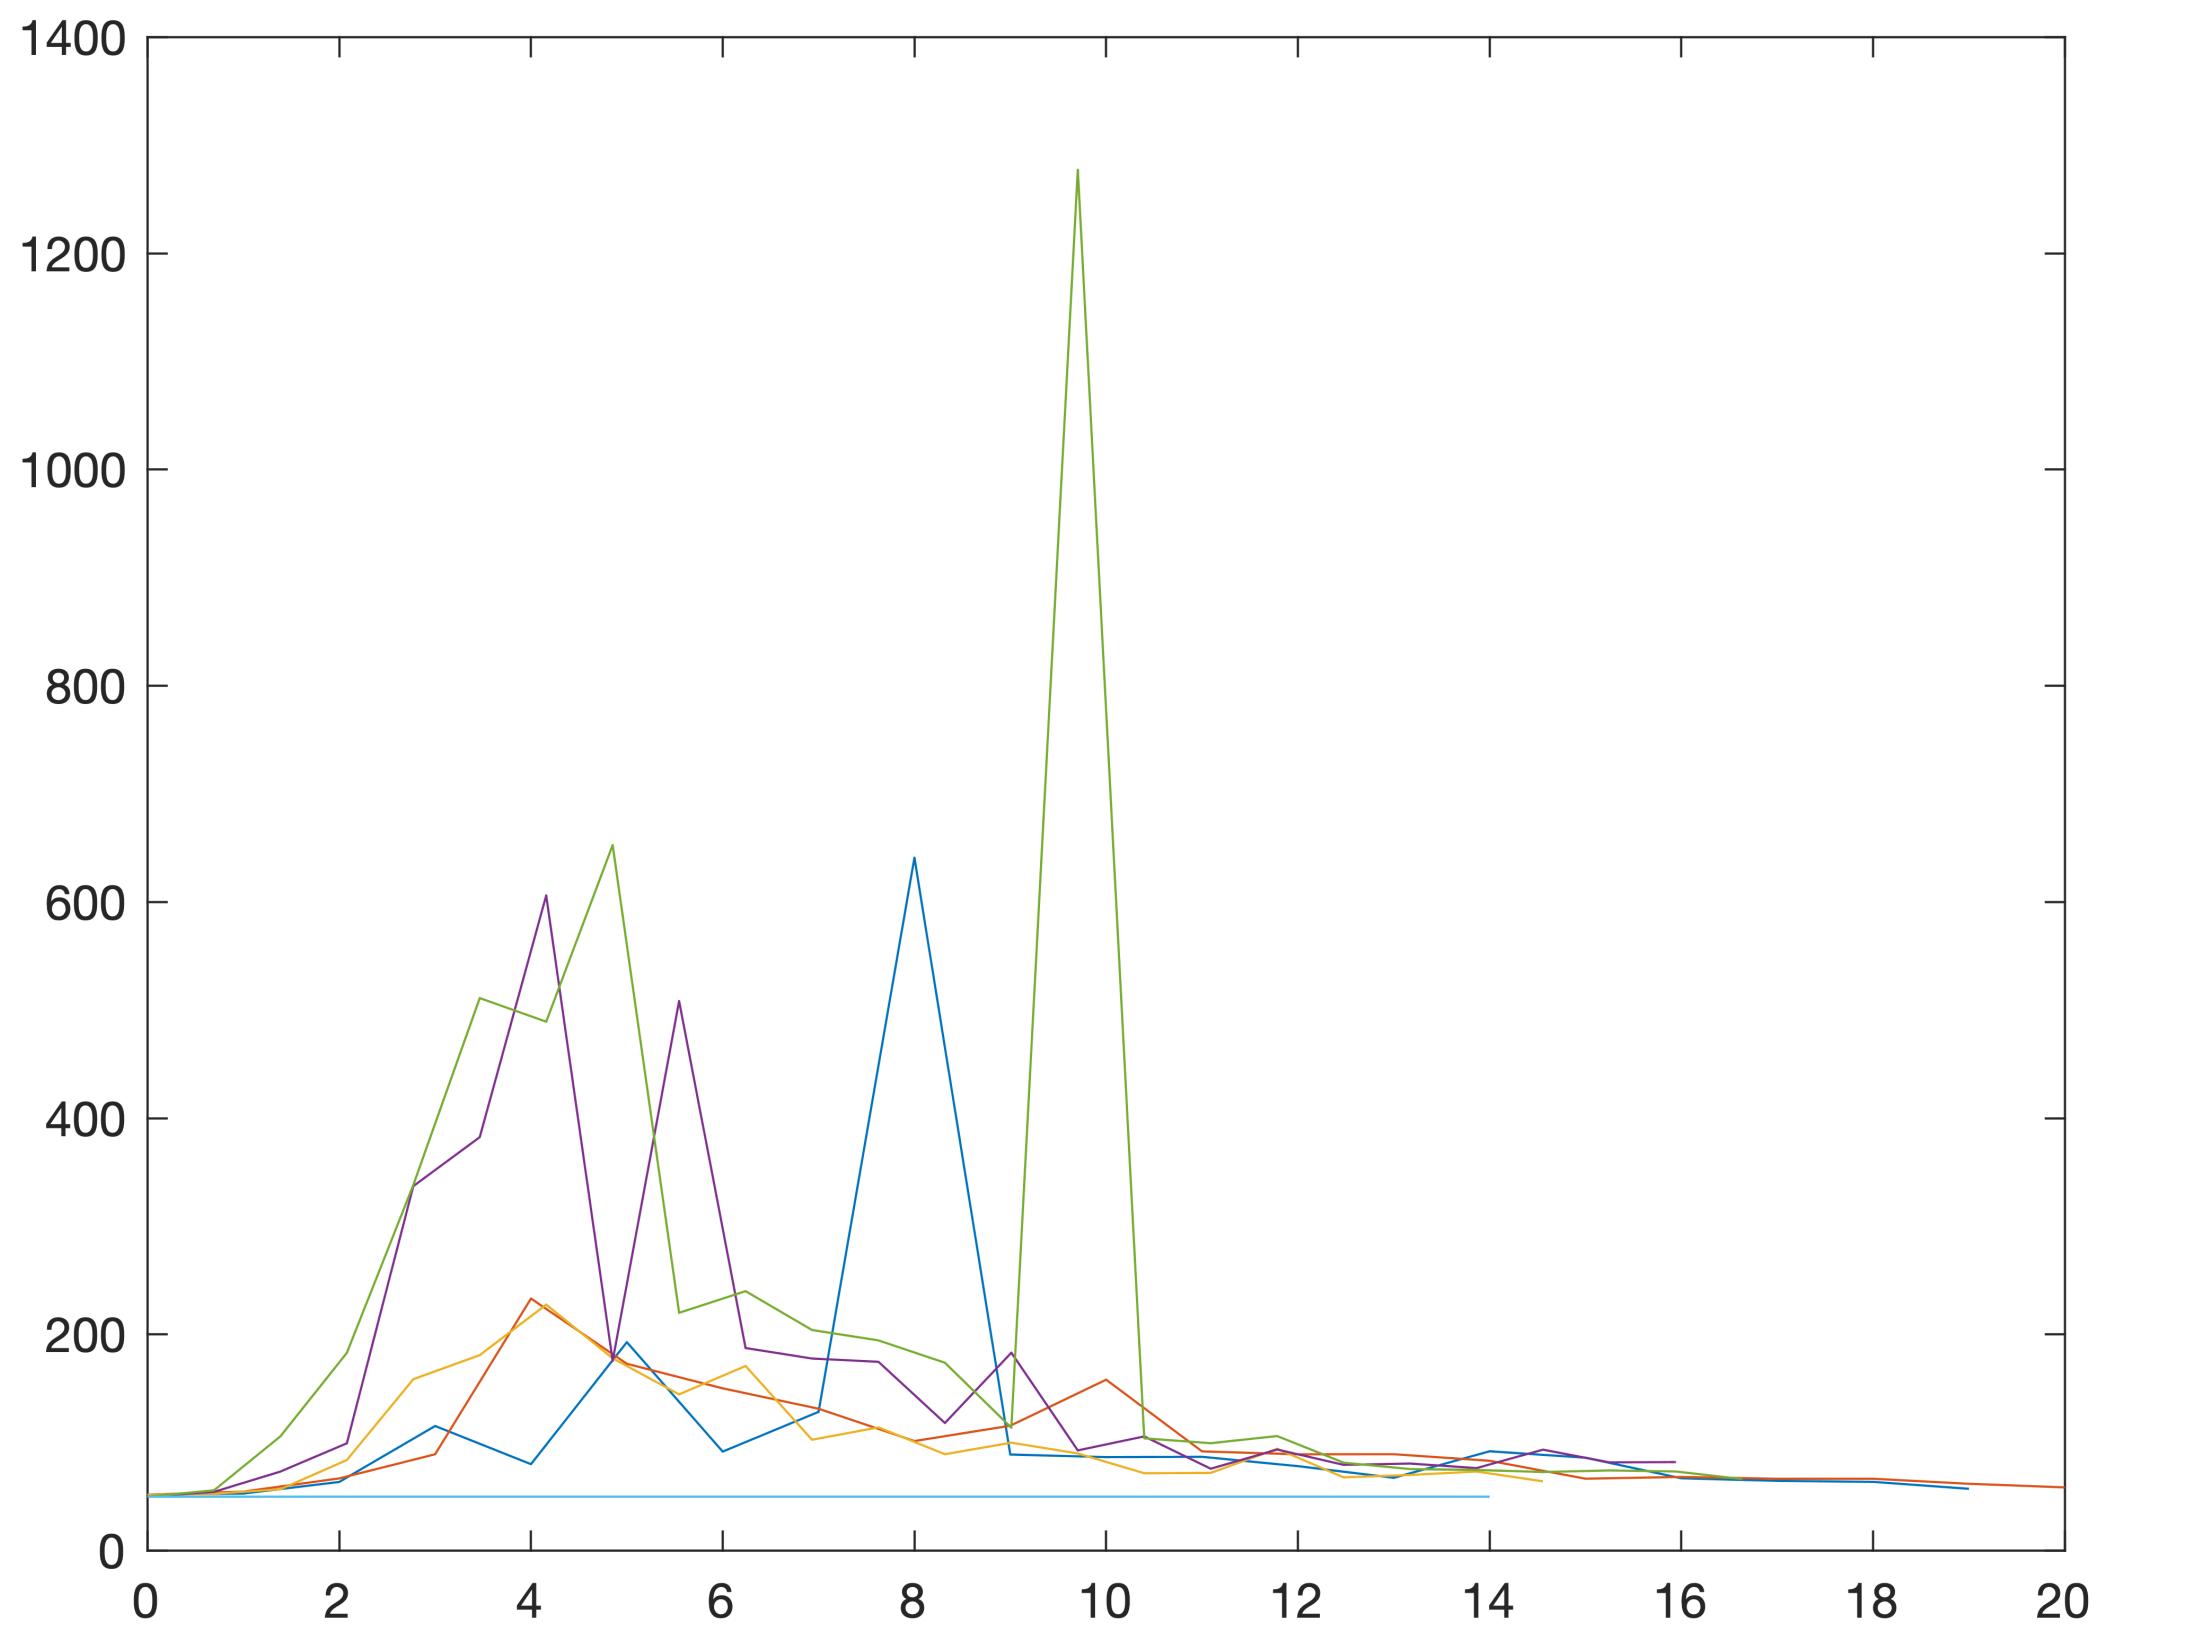


Result

Benchmark_S

- benchmark_s.c
 - *build the main loop*
 - *count the time*
 - *write the result*
- benchmark_n.sh
 - *change N & S automatically*

- **S+=S**
- **log10(s) - time**
- Too many noise !
- Hard to identify the s value

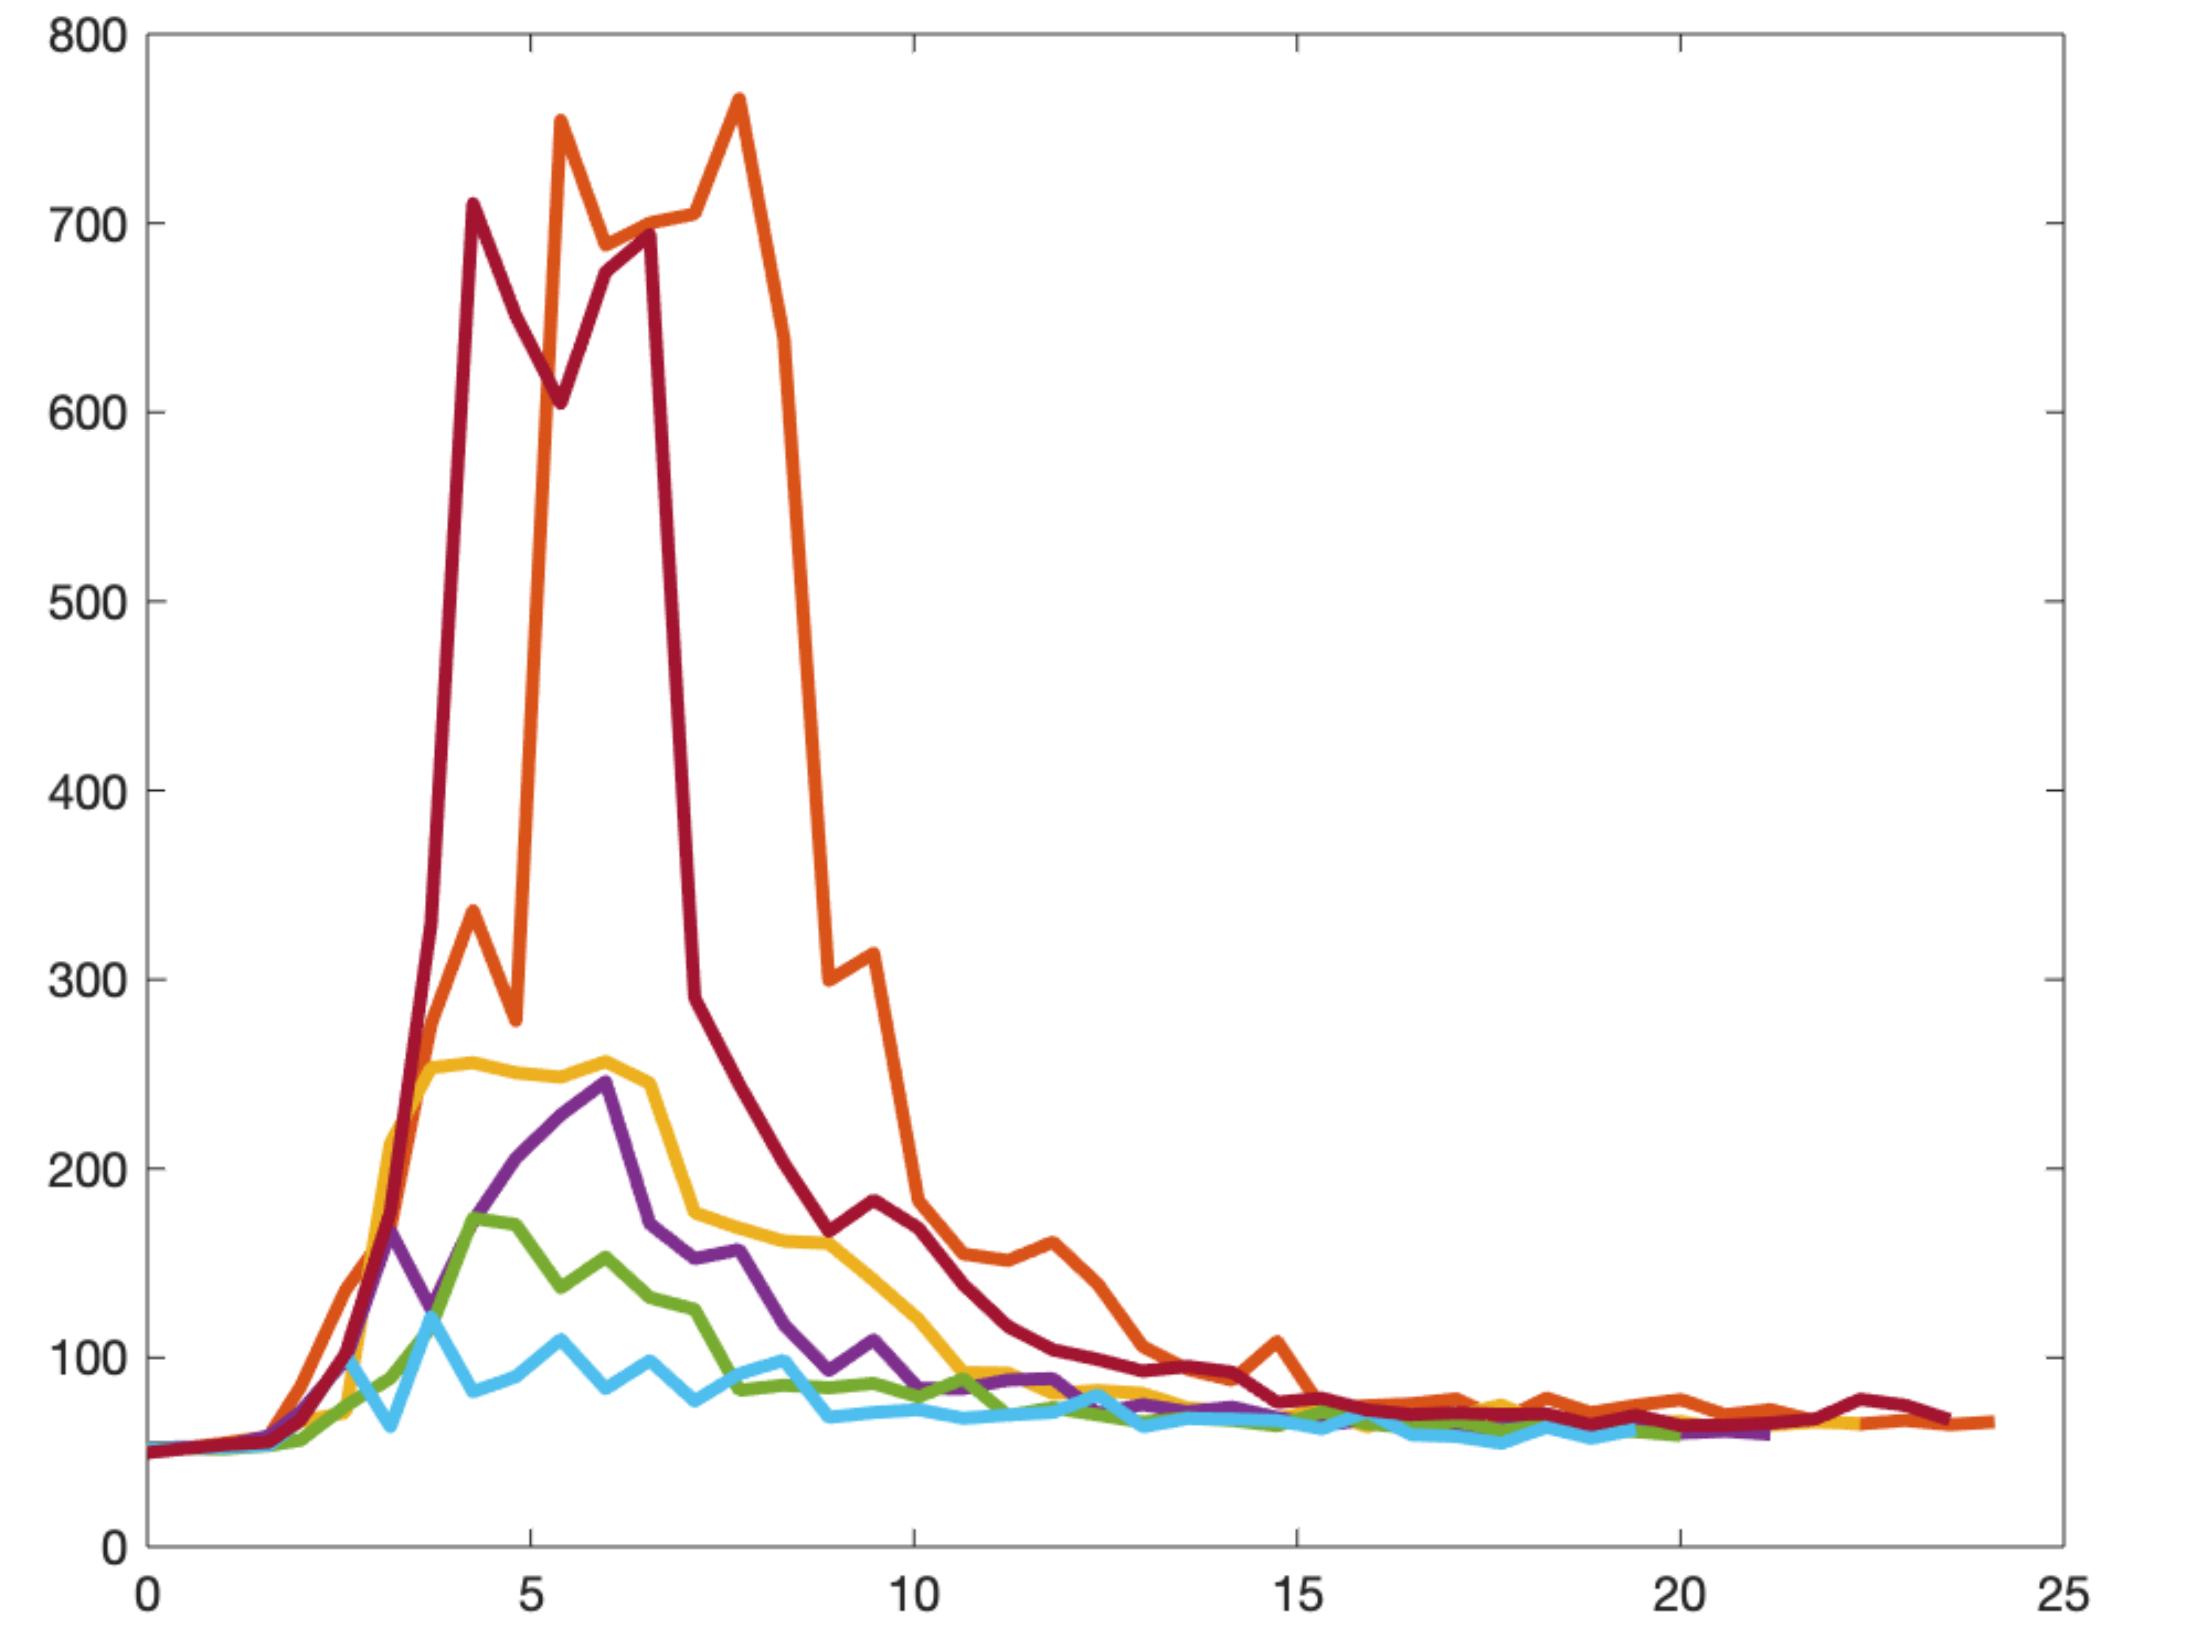


Result

Benchmark_S

- benchmark_s.c
 - *build the main loop*
 - *count the time*
 - *write the result*
- benchmark_n.sh
 - *change N & S automatically*

- $s+=s/2$
- $\log_2(s)$ - time
- Remove outlier
- Beautiful~

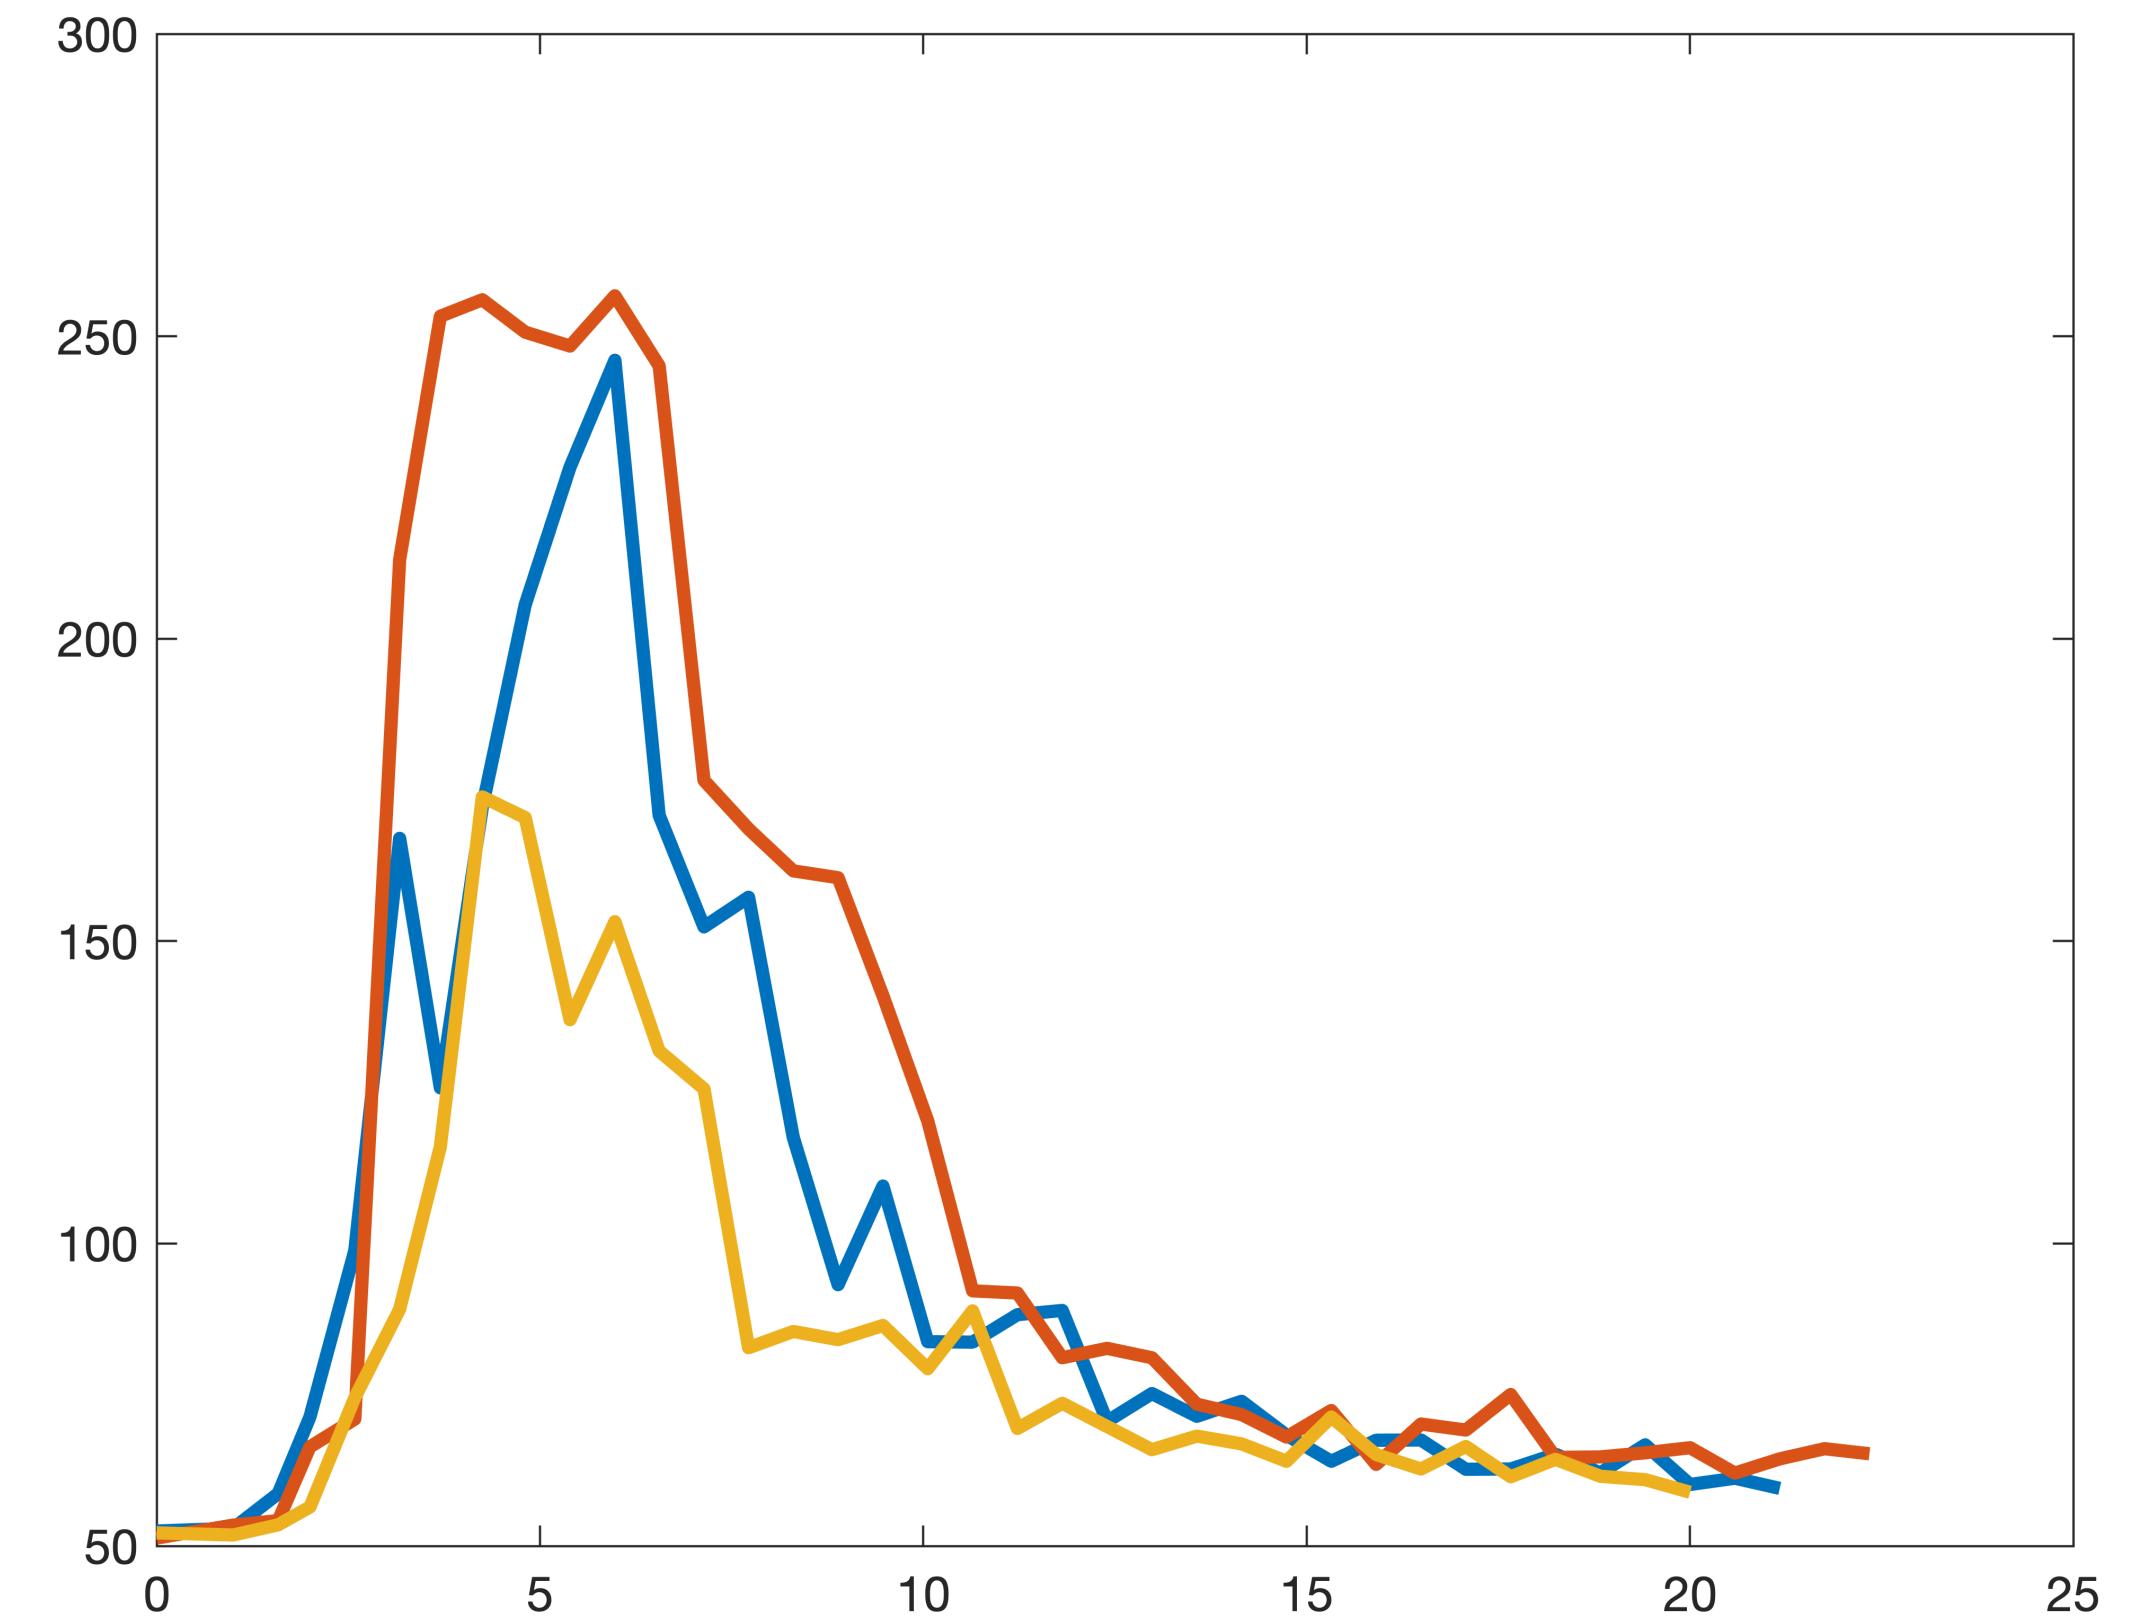


Result

Cache Line Size b

- Regime 2.a
- Line size b is the s value when the time value goes to maximum (under a relative small N)

- $N = 2/4/8 D$
- $\log_2(s) - \text{time}$
- $B = 32 * 4\text{byte} = 128 \text{ bytes}$

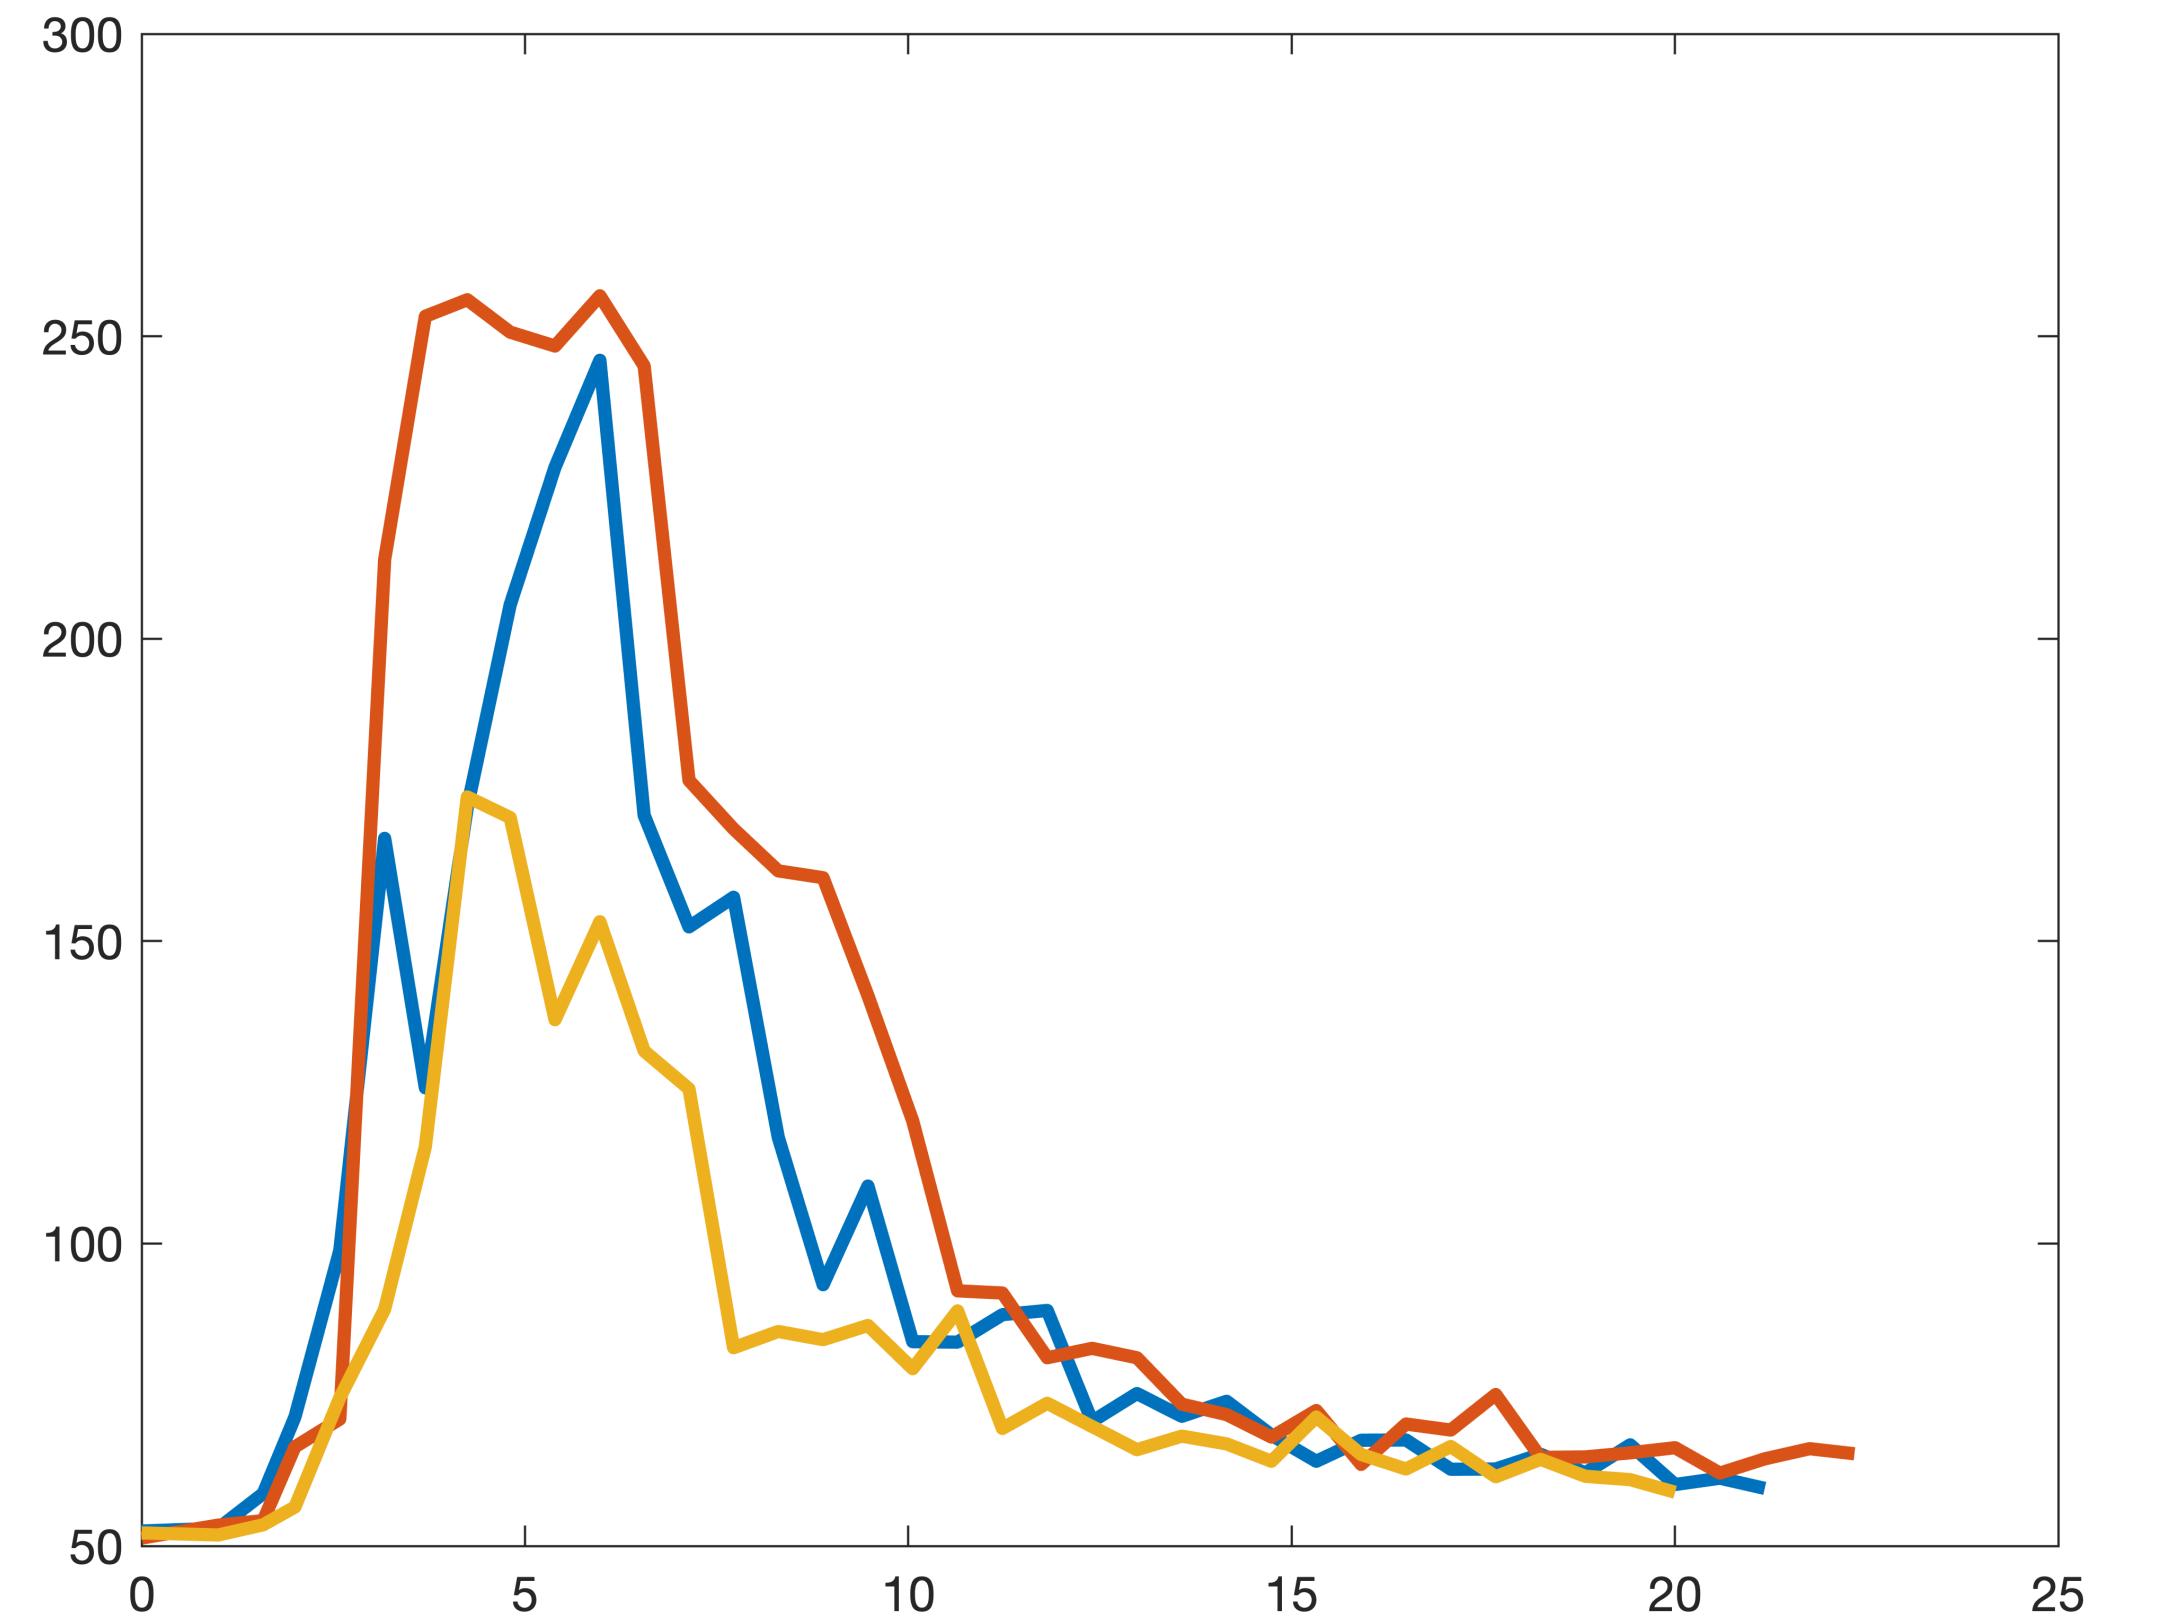


Result

Cache Associativity a

- Regime 2.b
- Cache associativity b is the N/s value when the time value drop to minimum (under a relative small N)

- $N = 2/4/8 D$
- $\log_2(s)$ - time
- $a = N/s_a$
 $= 4$

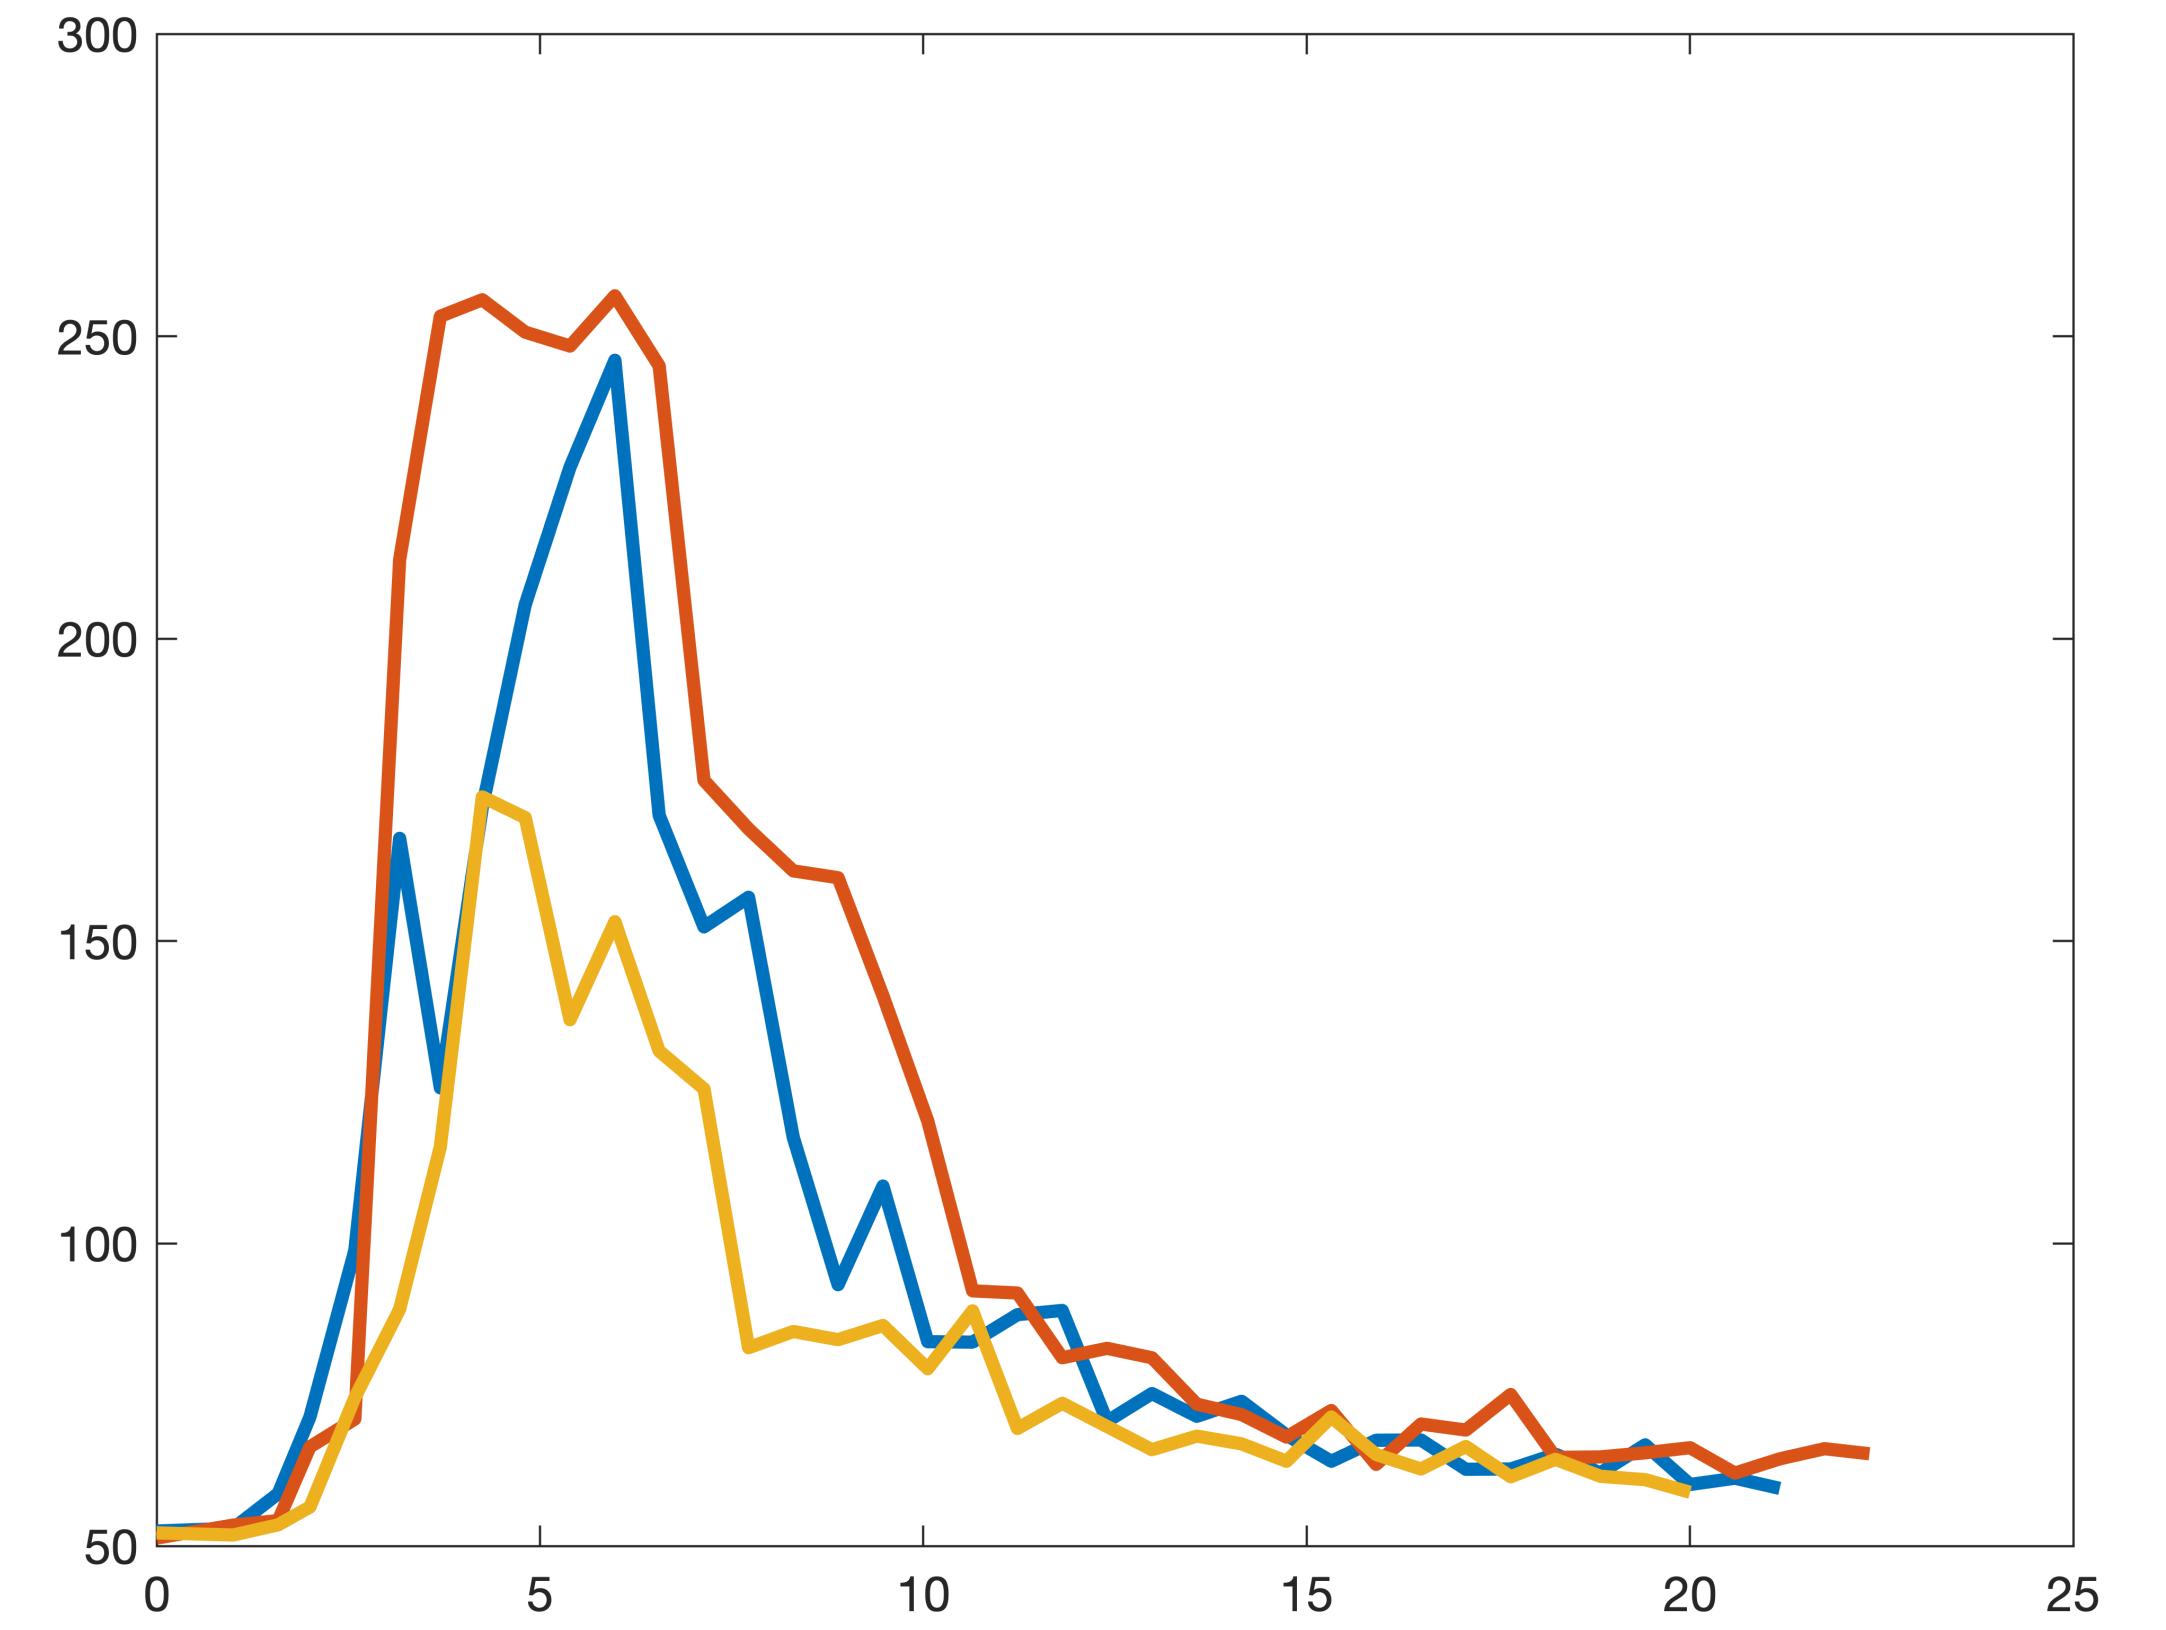


Result

Cache Miss penalty M

- Regime 2.a
- M is the time_n - time_o value when the time value goes to maximum (under a relative small N)

- $N = 2/4/8 D$
- $\log_2(s) - \text{time}$
- $M = 200 - 70 = 130 \text{ ns}$

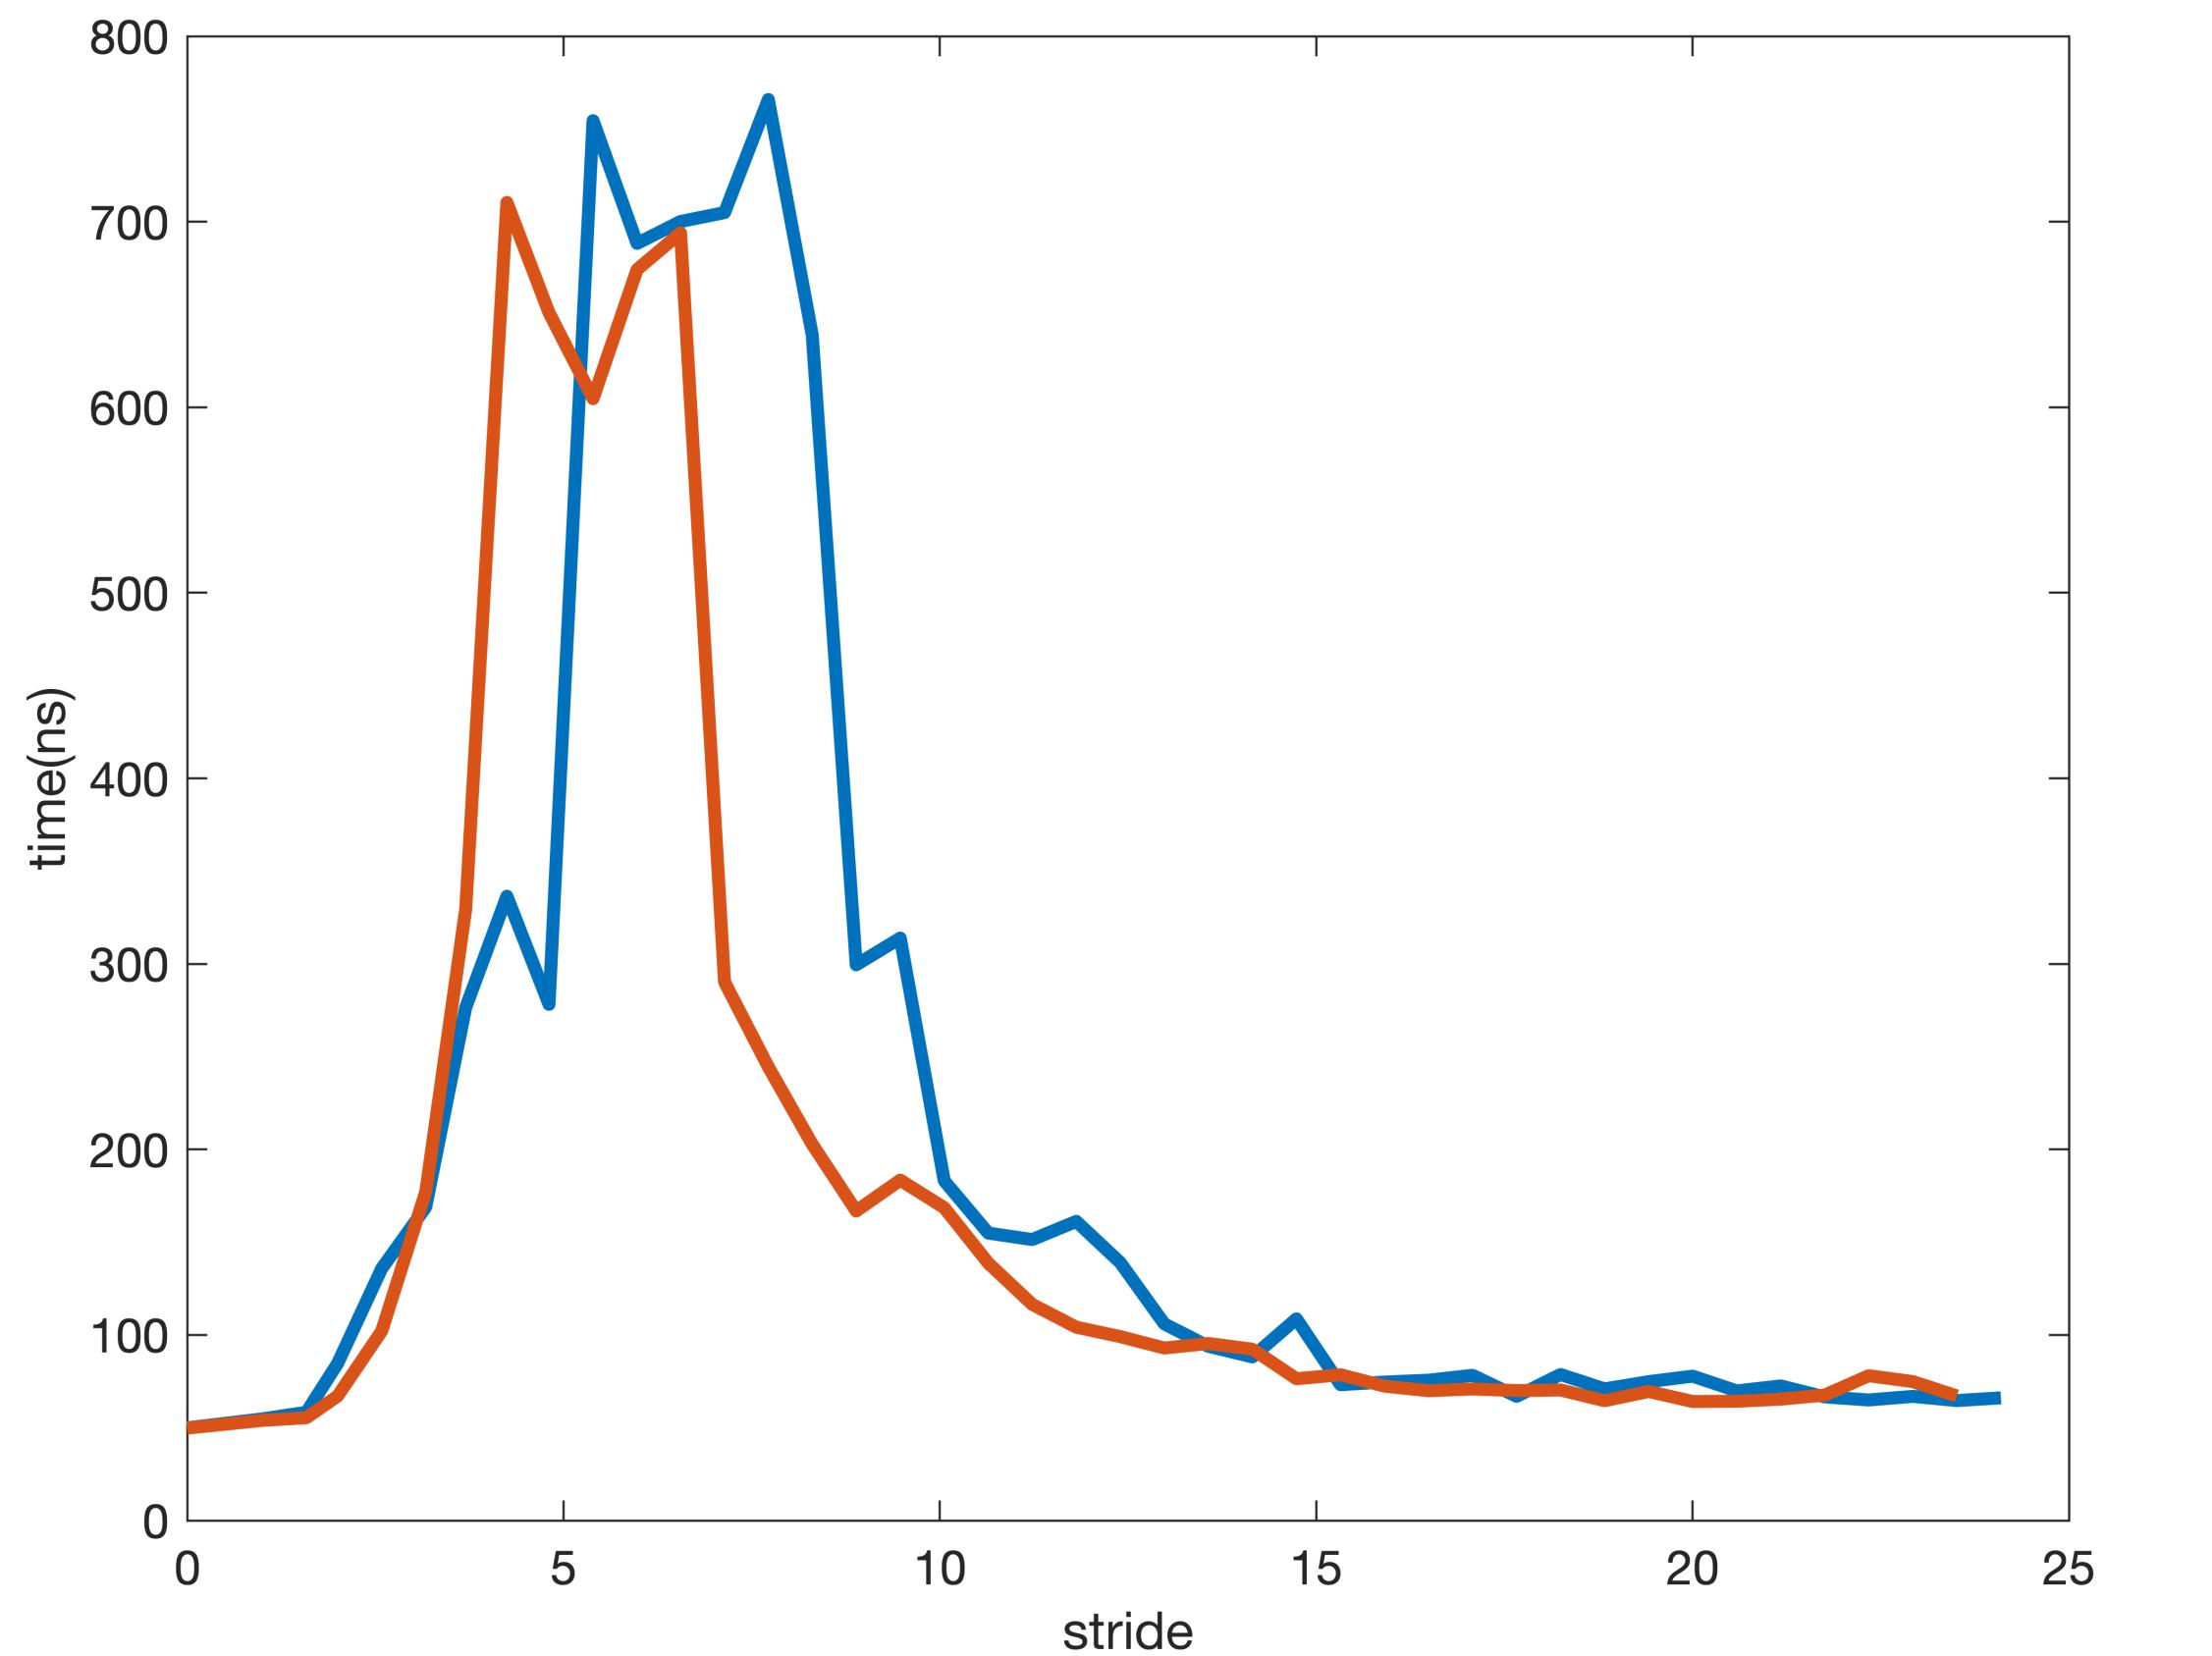


Result

Page Size p

- Regime 2.a
- Page size p is the s value when the time value goes to maximum (under a relative large N)

- $N = 16/32 D$
- $\log_2(s) - \text{time}$
- $p = 1024 * 4\text{byte} = 4K \text{ bytes}$



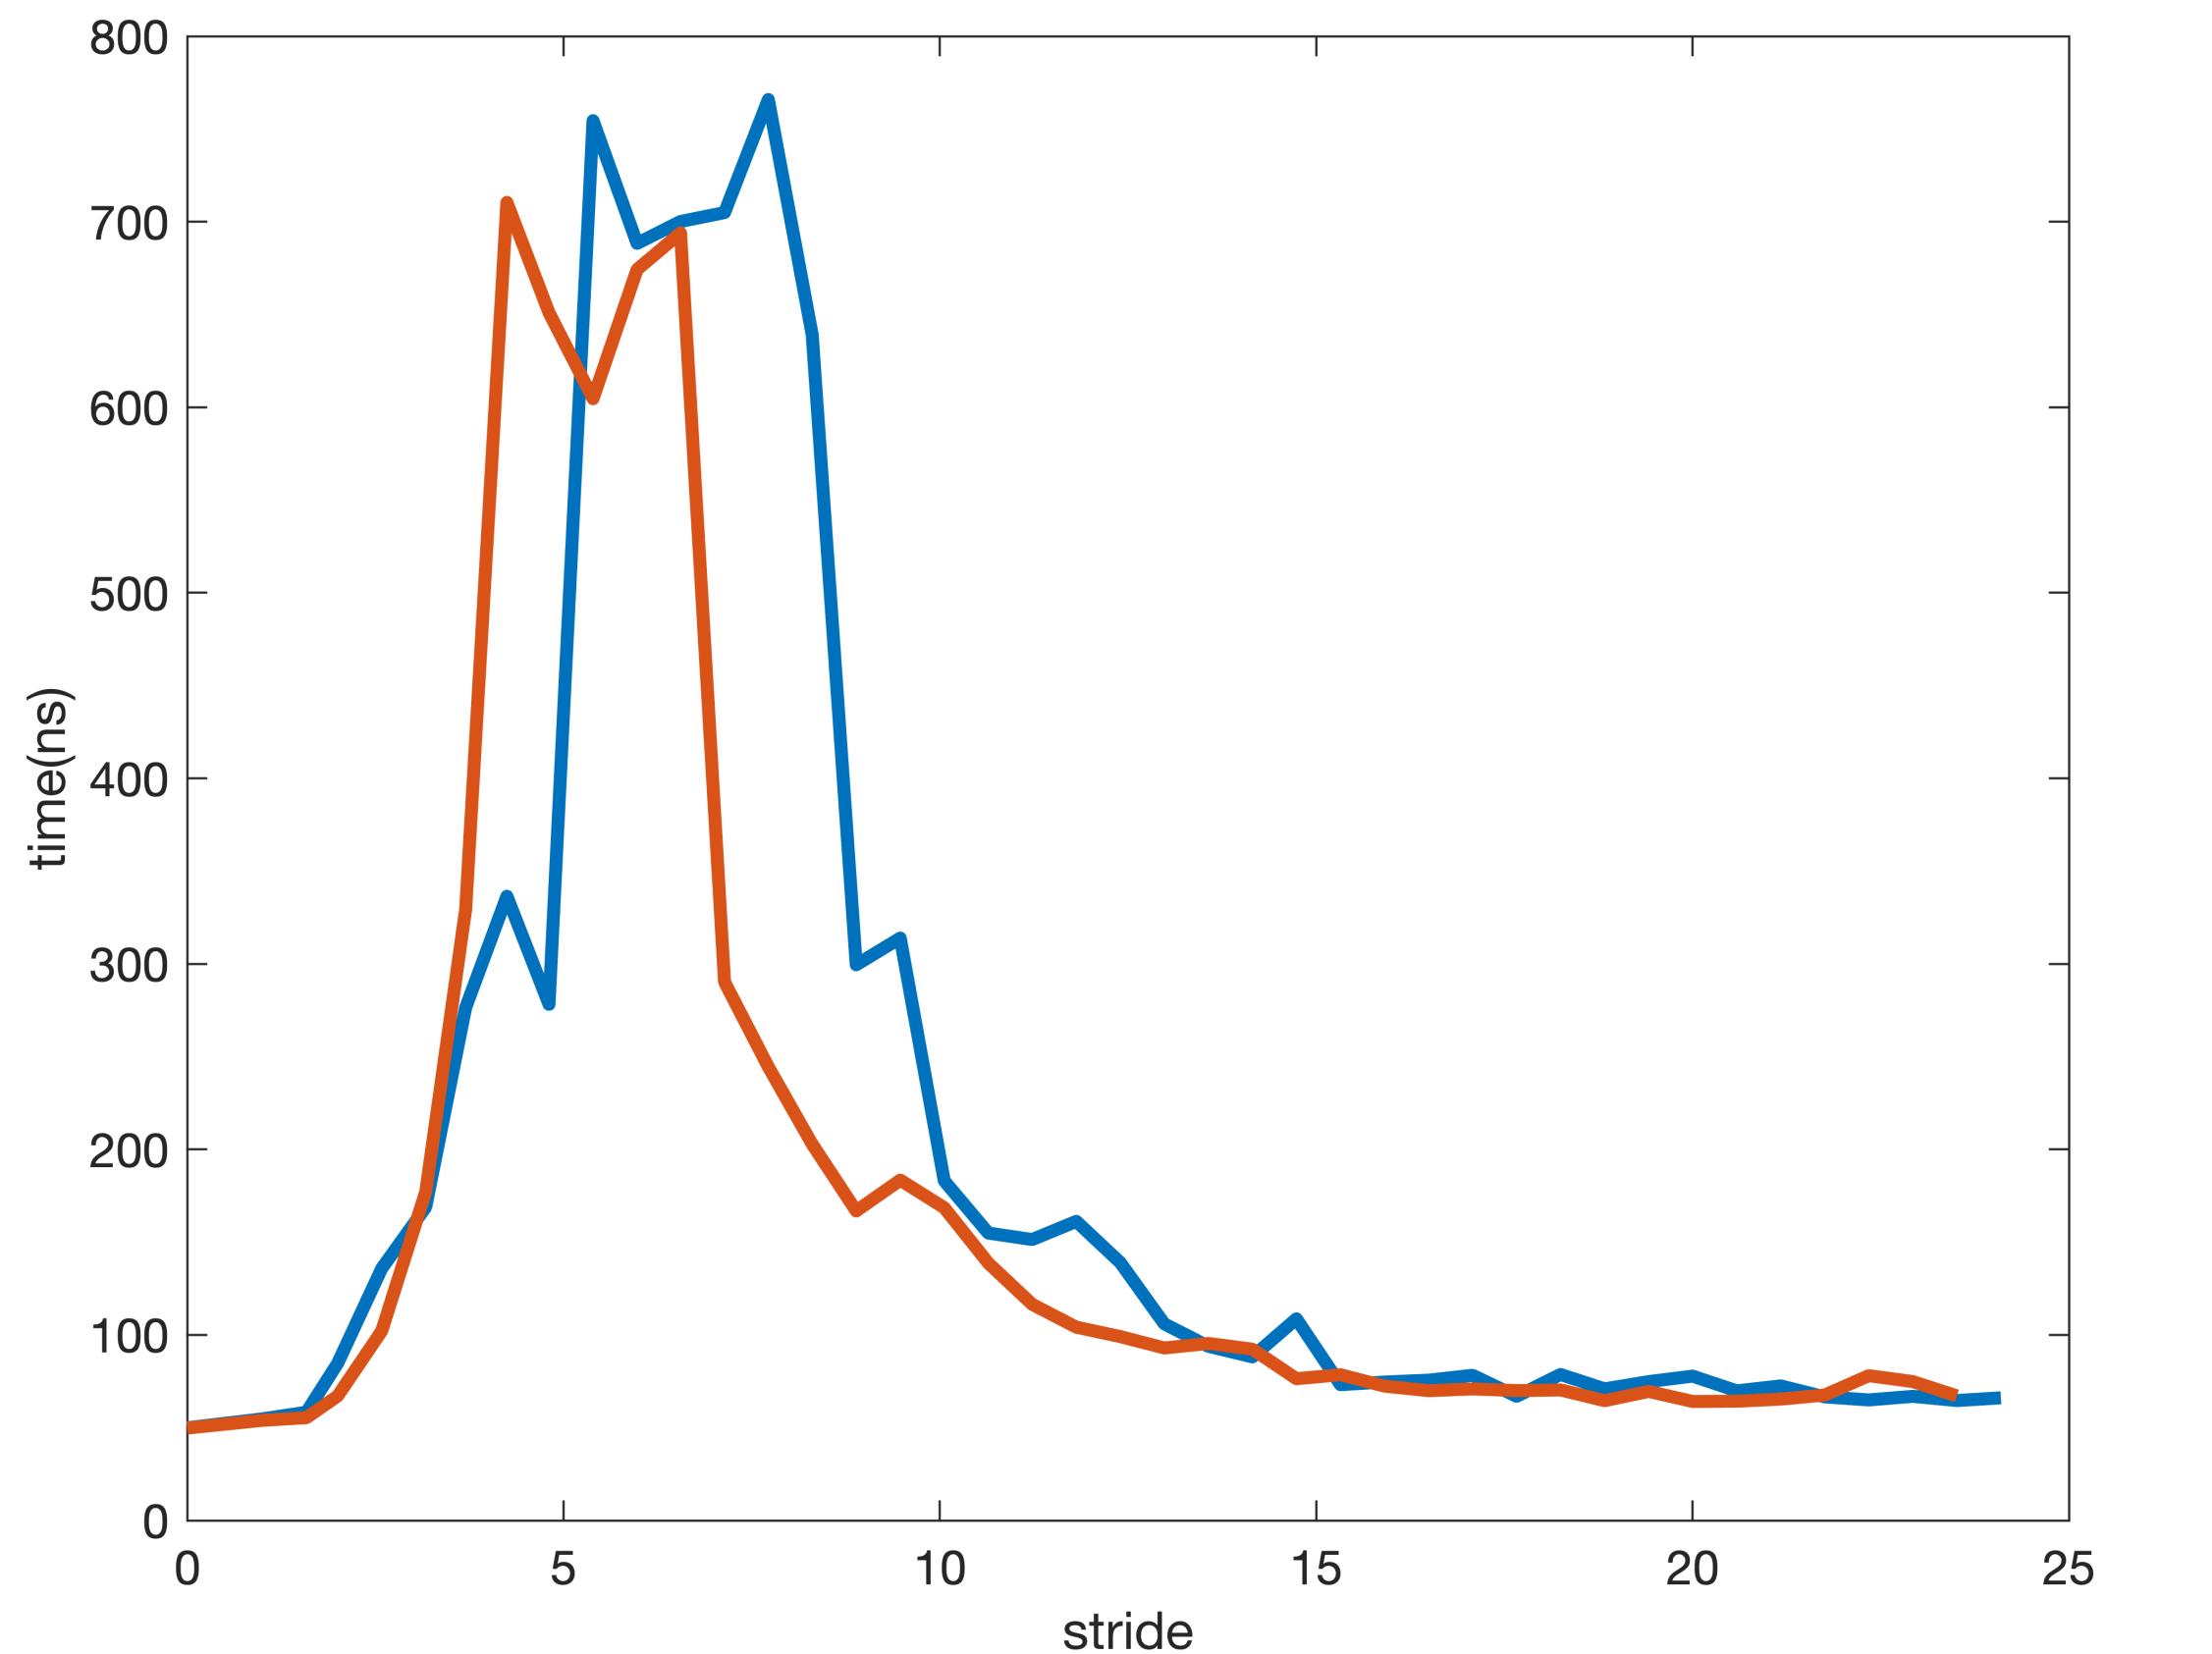
Result

TLB Associativity tlb_a

- Regime 2.b
- TLB associativity tlb_a is the N/s value when the time value drop to minimum

(under a relative large N)

- $N = 1/2/4/8 D$
- $\log_2(s) - \text{time}$
- $\text{tlb_a} = N/s_a = 4$

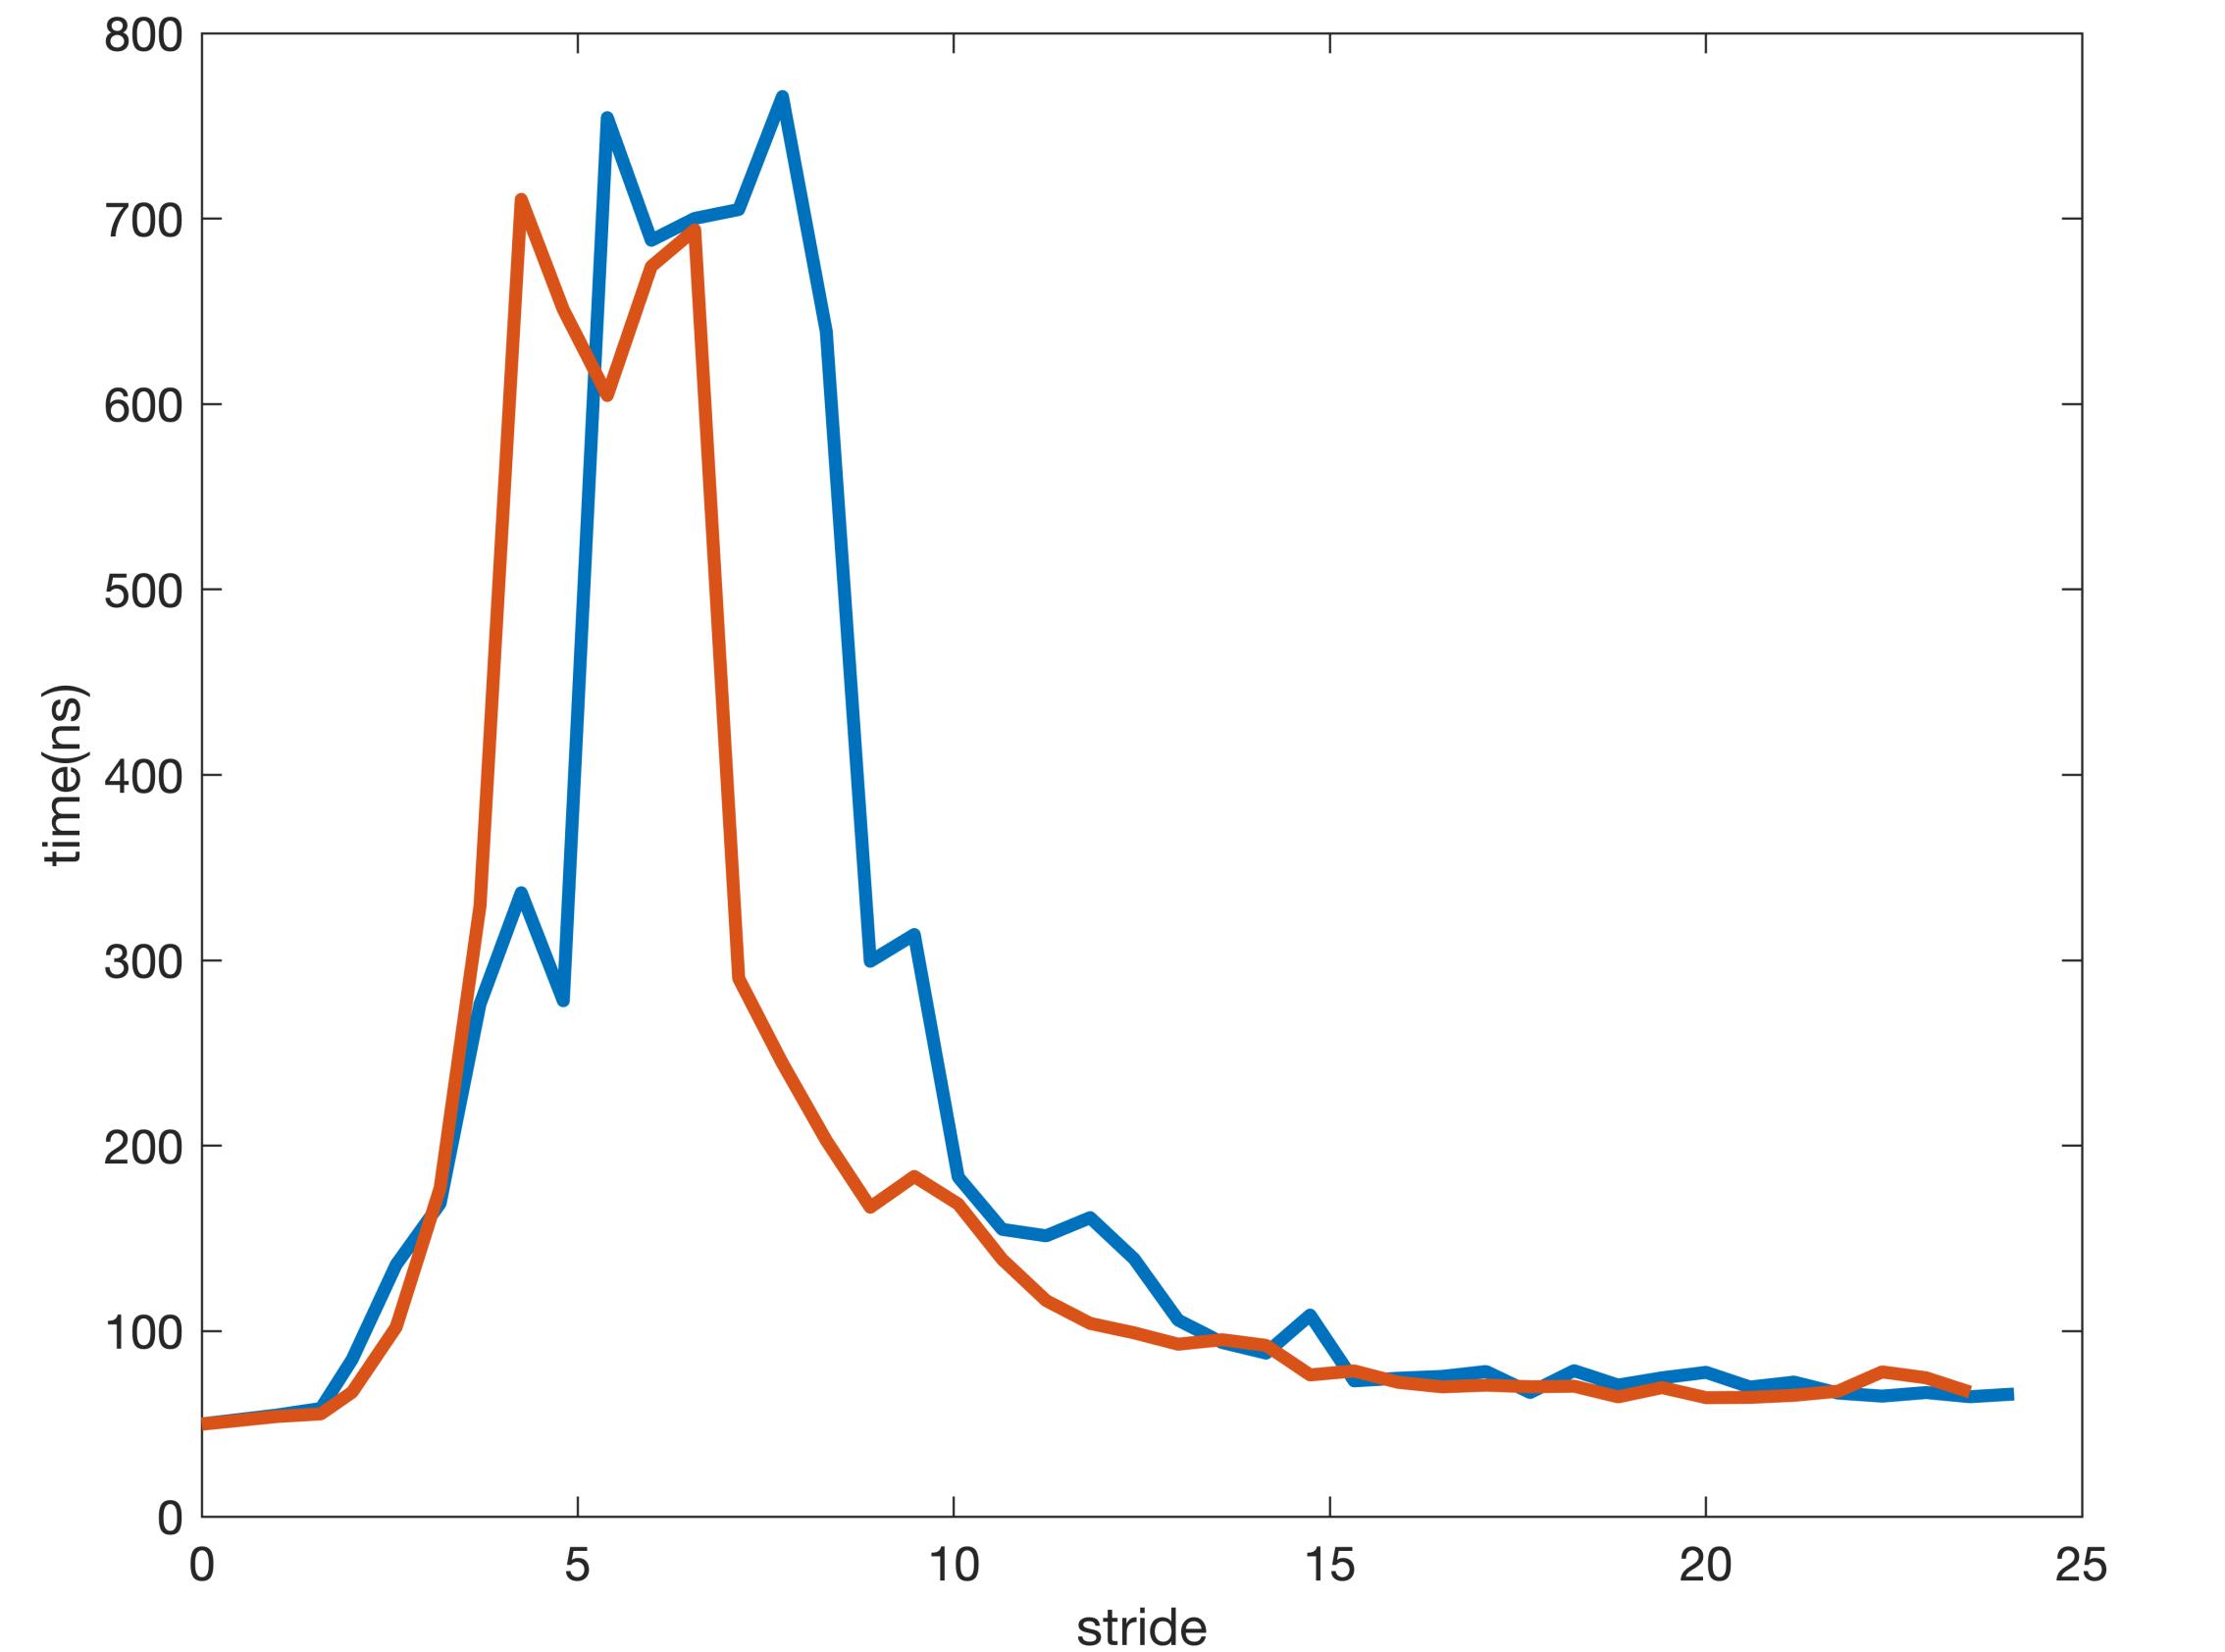


Result

TLB Miss penalty tlb_M

- Regime 2.a
- tlb_M is the time_n' - time_n value when the time value goes to maximum (under a relative smaller N)

- $N = 1/2/4/8 D$
- $\log_2(s) - \text{time}$
- $M = 700 - 130$
 $= 570 \text{ ns}$



Result

Conclusion

- I want to know the cpu's....
- Check the result

	<i>Cache Parameters</i>	<i>TLB Parameters</i>
<i>cache size</i>	6062KB	4KB
<i>associativity</i>	4	4
<i>line size</i>	128bytes	128bytes
<i>miss penalty</i>	130ns	570ns

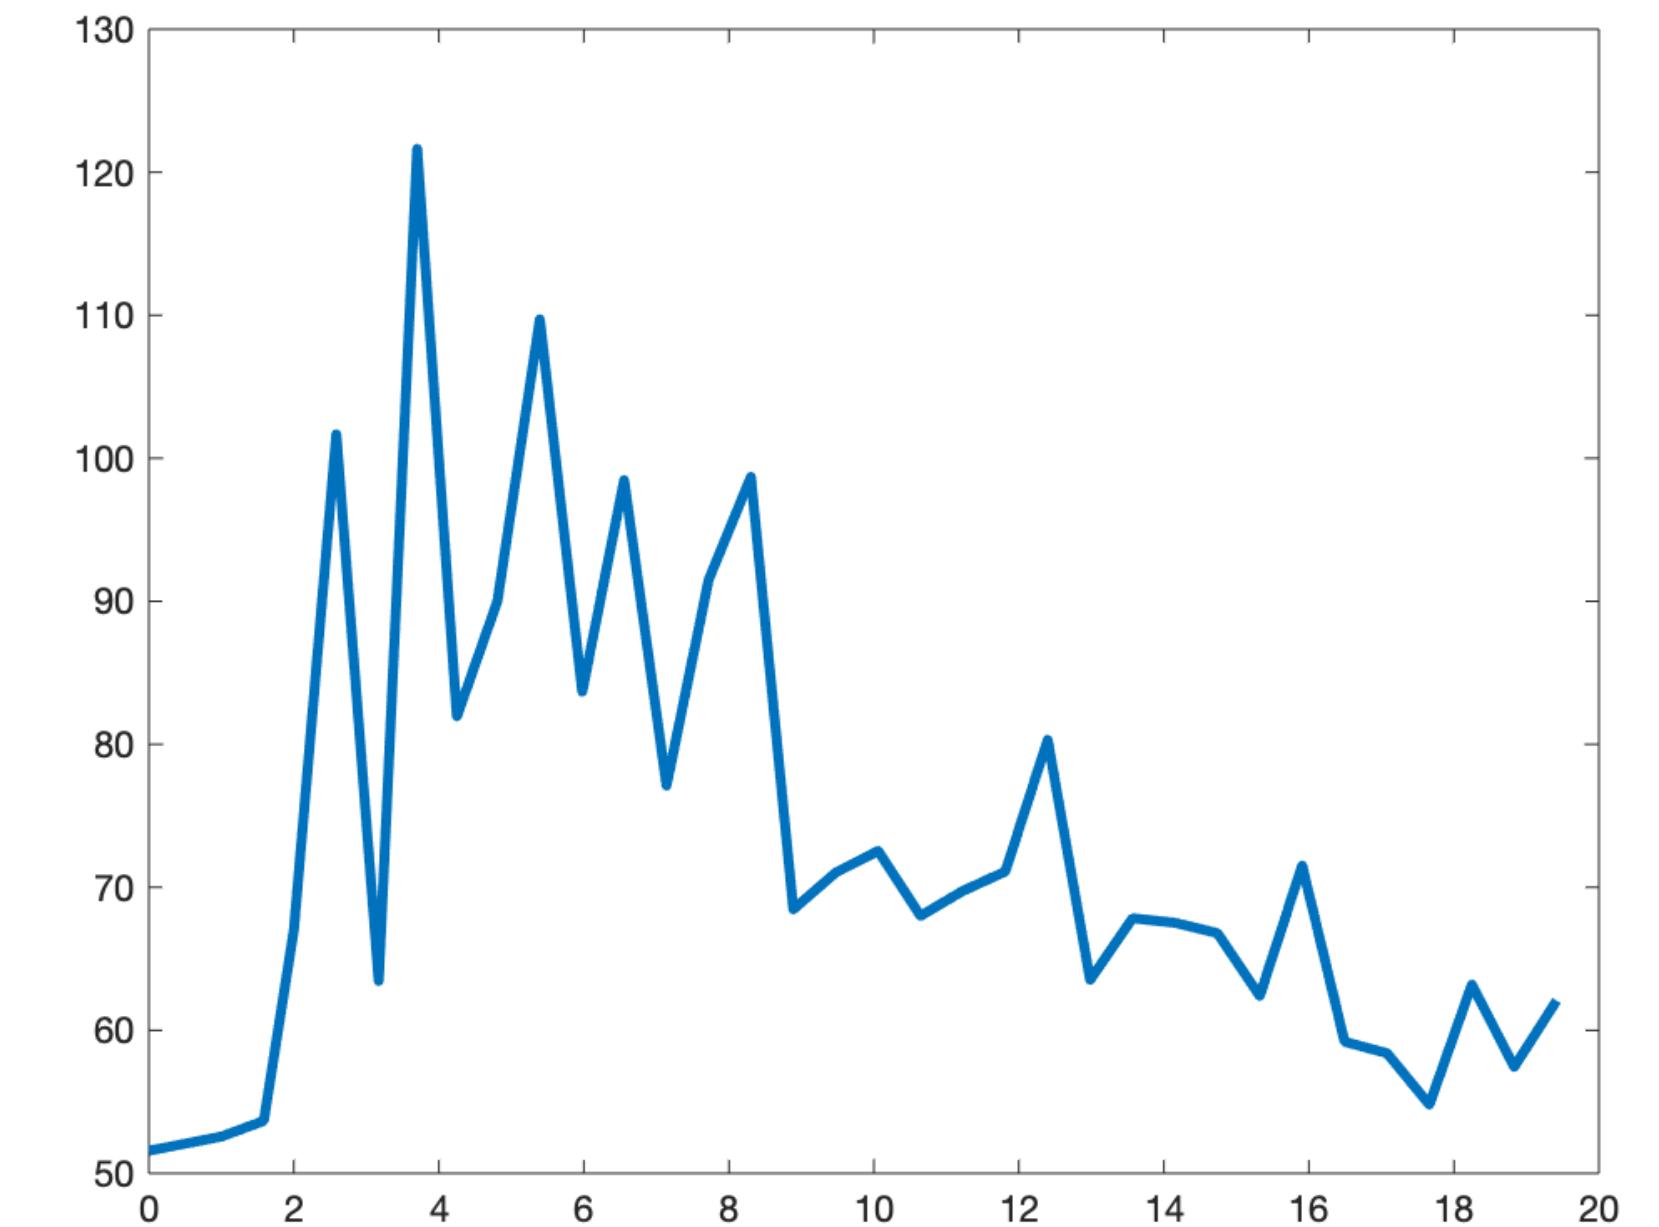
```
(base) yhc@yhc-Precision-7920-Tower:~/Desktop/os/cachesize$ cpuid | grep -i tlb
cache and TLB information (2):
0x63: data TLB: 2M/4M pages, 4-way, 32 entries
      data TLB: 1G pages, 4-way, 4 entries
0x03: data TLB: 4K pages, 4-way, 64 entries
0x76: instruction TLB: 2M/4M pages, fully, 8 entries
0xb6: instruction TLB: 4K, 8-way, 128 entries
0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
L1 TLB/cache information: 2M/4M pages & L1 TLB (0x80000005/eax):
L1 TLB/cache information: 4K pages & L1 TLB (0x80000005/ebx):
L2 TLB/cache information: 2M/4M pages & L2 TLB (0x80000006/eax):
L2 TLB/cache information: 4K pages & L2 TLB (0x80000006/ebx):
```

```
(base) yhc@yhc-Precision-7920-Tower:~/Desktop/os/cachesize$ cat /proc/cpuinfo
processor       : 0
vendor_id       : GenuineIntel
cpu family     : 6
model          : 85
model name     : Intel(R) Xeon(R) Bronze 3104 CPU @ 1.70GHz
stepping        : 4
microcode       : 0x2006b06
cpu MHz         : 1700.000
cache size      : 8448 KB
physical id    : 0
siblings        : 6
core id         : 0
cpu cores       : 6
```

Problems.....

Fight with the 'prefetch'

- Prefetch: fetch more than you want
- In our case:
 - Set stride=1 and get very strange result
- My solution:
 - Use stride=1000(100 may work as well)



Problems.....

How to measure the time

- Which api to use
- `clock()`: only ms level
- My solution: search online
 - A good api: `clock_gettime`

- Which measurement to use
- Count average time of each loop
 - Failed when N went bigger.
- Count the time of each iteration
 - Cost a lot.
- My solution: spend some time.....

实验体会

- 实验难度不小
 - 论文非常清晰，但是复现起来困难重重
 - Debug很痛苦，尤其是预取问题刚出现时
- 能力提升很大
 - 理论-实现的过程，非常有趣
 - 锻炼分析问题的思维
 - 感谢老师 ~

Thanks For Listening!

- Github id: yyyyhc